

Thread Safety

Race Conditions and Thread Synchronization (Locking)

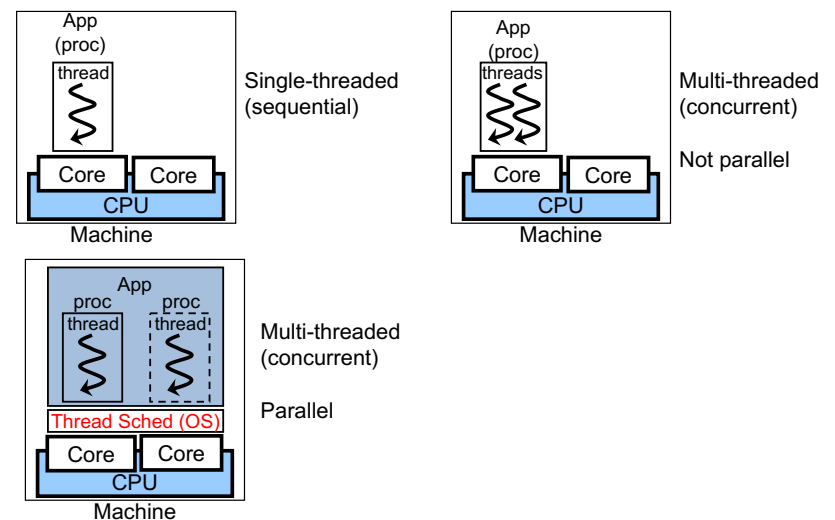
- Threads are a powerful tool to make your code more responsive and efficient.
- However, multi-threaded code can raise **“thread safety” issues** if it is poorly written.
- 2 major thread safety issues
 - Race conditions (data races)
 - Mess up the consistency of data shared among threads
 - Deadlock
 - Make code execution stuck.
- **Thread-safe code** is free from those 2 issues.

Race Conditions (Data Races)

- Threads run **independently**.
 - No coordination among threads by default.
 - c.f. MCTest, PrimeNumberGenerator
 - join() allows threads to coordinate with each other.
- They can **share** objects/data.
 - Exception: Local variables are NOT shared among threads.
- They can **mess up** the consistency of the shared objects/data.
 - A thread can write some data to a variable when another thread is reading data from the variable.
 - A thread can write some data to a variable when another thread is writing different data to the variable.

3


Single- and Multi-threaded Programs



4


In CS681...

- We always assume a single CPU core that run multiple threads.
 - The most conservative scenario.
- If your code is thread-safe in the most conservative scenario, it is always thread-safe in less conservative scenarios as well.



An Example Race Condition:

ThreadUnsafeBankAccount.java



- The variable “balance” is shared by 2 threads.
 - They access the variable independently.
- ```


public void deposit(double amount){
 System.out.print("Current balance (d): " + balance);
 double newBalance = balance + amount;
 System.out.println(", New balance (d): " + newBalance);
 balance = newBalance;
}

public void withdraw(double amount){
 System.out.print("Current balance (w): " + balance);
 double newBalance = balance - amount;
 System.out.println(", New balance (w): " + newBalance);
 balance = newBalance;
}

```

6

## How Can This Happen?



### Desirable output:

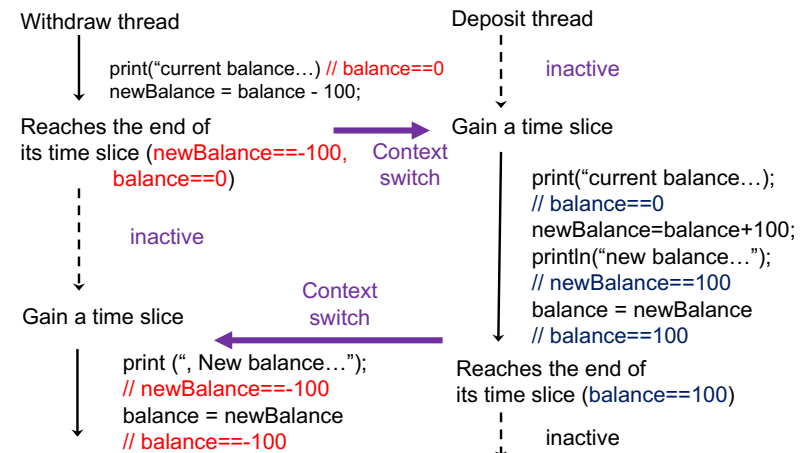
- Current balance (w): 0.0, New balance (w): -100.0
- Current balance (d): -100.0, New balance (d): 0.0
- Current balance (w): 0.0, New balance (w): -100.0
- Current balance (d): -100.0, New balance (d): 0.0
- Current balance (d): 0.0, New balance (d): 100.0
- Current balance (w): 100.0, New balance (w): 0.0
- ...

### In reality:

- Current balance (w): 0.0, Current balance (d): 0.0, New balance (d): 100.0
- , New balance (w): -100.0

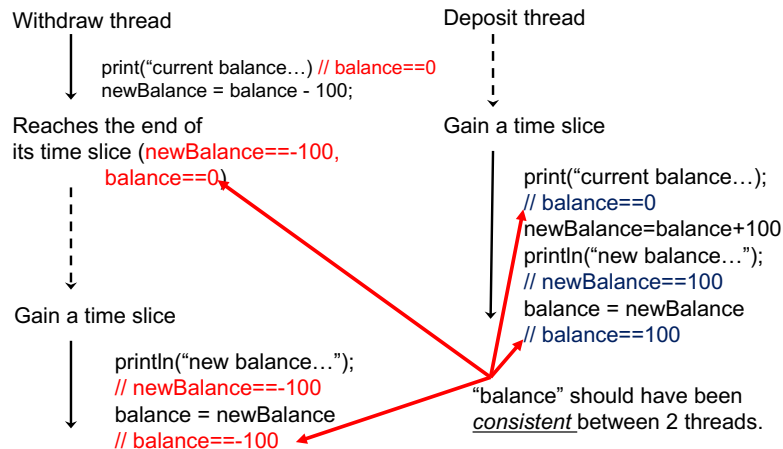
7

- Current balance (w): 0.0, Current balance (d): 0.0, New balance (d): 100.0
- , New balance (w): -100.0



8

- Current balance (w): 0.0 Current balance (d): 0.0, New balance (d): 100.0
- , New balance (w): -100.0



9

## The Source of the Problem: Visibility


- ThreadUnsafeBankAccount.java is **NOT thread safe**.
  - Race conditions can occur.
- Race conditions occur due to **visibility** issues.
  - The current (most up-to-date) value of the shared variable (e.g. "balance") is not visible for all threads.

## Race Conditions (Data Races)

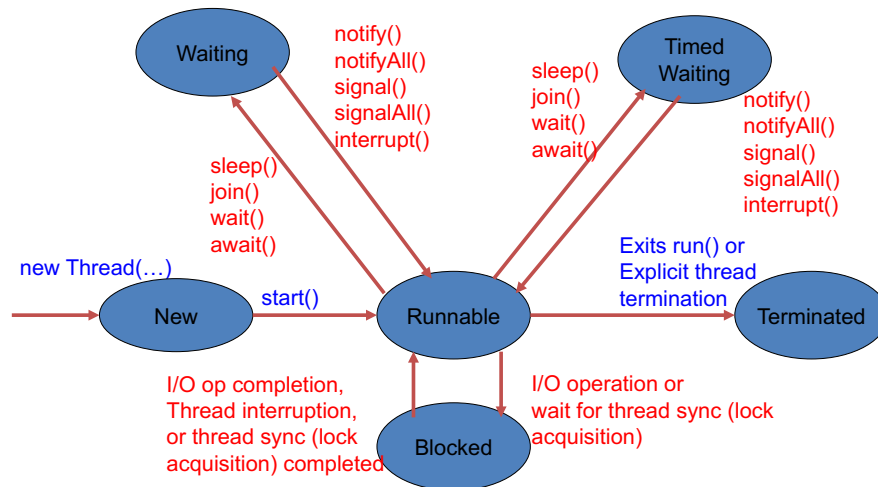
- All threads
  - Run in their **race** to complete their tasks.
  - Manipulate a shared object/data **independently**.
- The end result depends on which of them happens to win the race.
  - No guarantees on the order of thread execution.
  - No guarantees on how many tasks a thread can perform in a single CPU time slice.
  - No guarantees on the end result on shared data.

## Note: Thread States

- Both **"active"** and **"inactive"** threads are in the **Runnable** state.
  - The **Runnable** state does NOT distinguish if a thread is actively running on a CPU core or it is inactive waiting for its next turn.
  - The **Waiting** state does NOT mean that a thread is runnable but inactive.



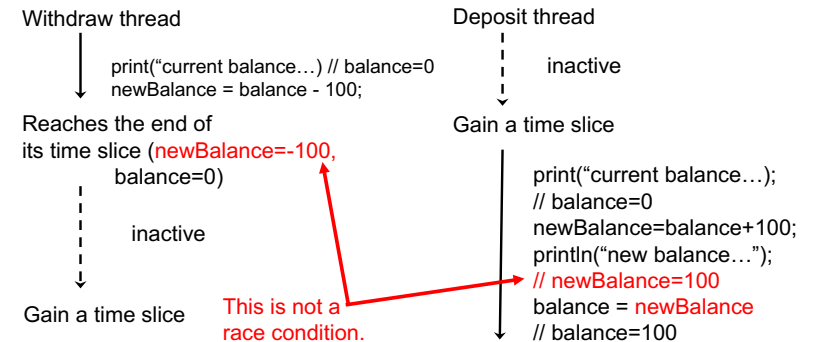
## States of a Thread



13

## Note: Local Variables

- Local variables are NOT shared by threads.
  - It is created and maintained in a **thread-by-thread** manner.
    - The “withdraw” thread has no access to a value of newBalance that the “deposit” thread has created.
    - The “deposit” thread has no access to a value of newBalance that the “withdraw” thread has created.



## Another Example:

### ThreadUnsafeBankAccount2

- Race conditions never occur due to local variables.
- Focus on non-local (i.e. **shared**) variables in debugging threaded code.
- Two other example local variables

```

- public void deposit(double amount){
 System.out.print("Current balance (d): " + balance);
 double newBalance = balance + amount;
 System.out.println(", New balance (d): " + newBalance);
 balance = newBalance;
}

- public void withdraw(double amount){
 System.out.print("Current balance (w): " + balance);
 double newBalance = balance - amount;
 System.out.println(", New balance (w): " + newBalance);
 balance = newBalance;
}

```

- Local variables (newBalance) are removed from ThreadUnsafeBankAccount

```

- public void deposit(double amount){
 balance = balance + amount;
}

- public void withdraw(double amount){
 balance = balance - amount;
}

```

- Output

```

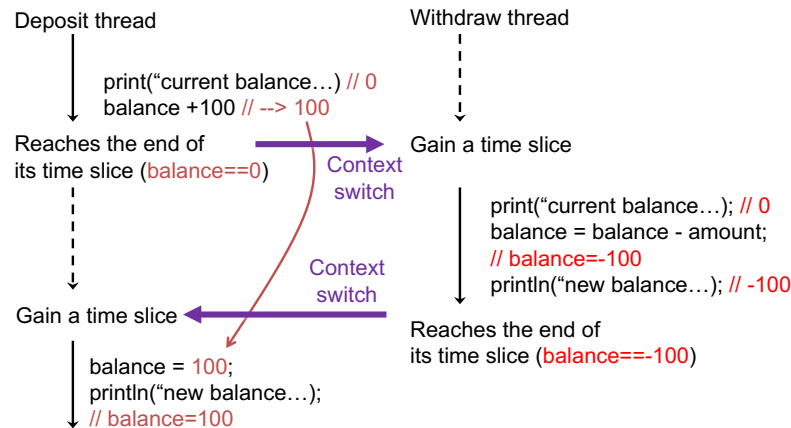
- Current balance (d): 0.0, New balance (d): 100.0
- Current balance (w): 100.0, New balance (w): 0.0
- Current balance (d): 0.0, New balance (d): 100.0
- Current balance (w): 100.0, New balance (w): 0.0
- Current balance (d): 0.0Current balance (w): 0.0, New balance (w): -100.0
- , New balance (d): 100.0

```

16

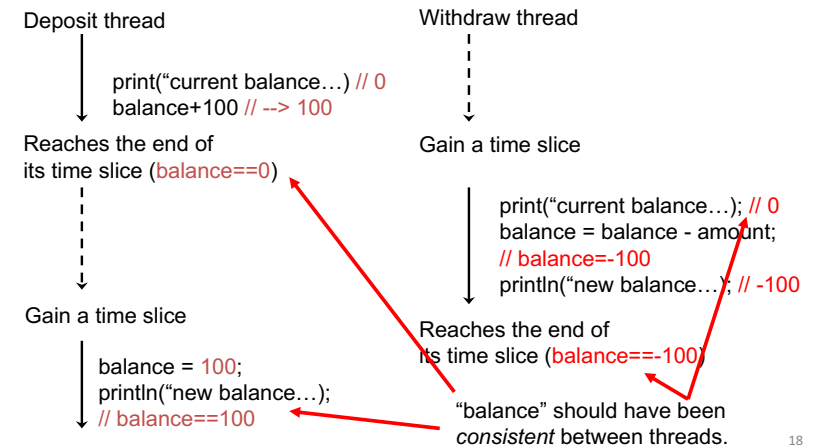
# How Can This Happen?

- Current balance (d): 0.0 Current balance (w): 0.0, New balance (w): -100.0
- , New balance (d): 100.0



17

- Current balance (d): 0.0 Current balance (w): 0.0, New balance (w): -100.0
- , New balance (d): 100.0



18

## Thread Synchronization

- This is not a solution: `balance -= amount;`
  - Just a **syntactic sugar** for `balance = balance - amount;`
- All threads
  - Run in their *race* to complete their tasks.
  - Manipulate a shared object/data independently.
- The end result depends on which of them happens to win the race.
  - No guarantees on the order of thread execution.
  - No guarantees on how many tasks a thread can perform in a single CPU time slice.
  - No guarantees on the end result on shared data.

19

- Need to **synchronize** threads
  - i.e., Need to **serialize** their concurrent access to a shared variable
- Thread synchronization enables **serialized** (or **atomic** or **exclusive**) access to a shared variable
  - Allows **only one thread** to access a shared variable at a time
    - Forces the other threads to wait and take turn to access it.
  - Prevents the "deposit" thread from depositing money to "balance" when the "withdraw" thread is withdrawing money from "balance."
  - Prevents the "withdraw" thread from withdrawing money from "balance" when the "deposit" thread from depositing money to "balance."

20

# Locks

- Thread synchronization allows programmers to write *atomic code* (a.k.a critical section).
  - When a thread is running it, no other threads can run it.
  - No intermediate result/state can be revealed/exposed to other threads.

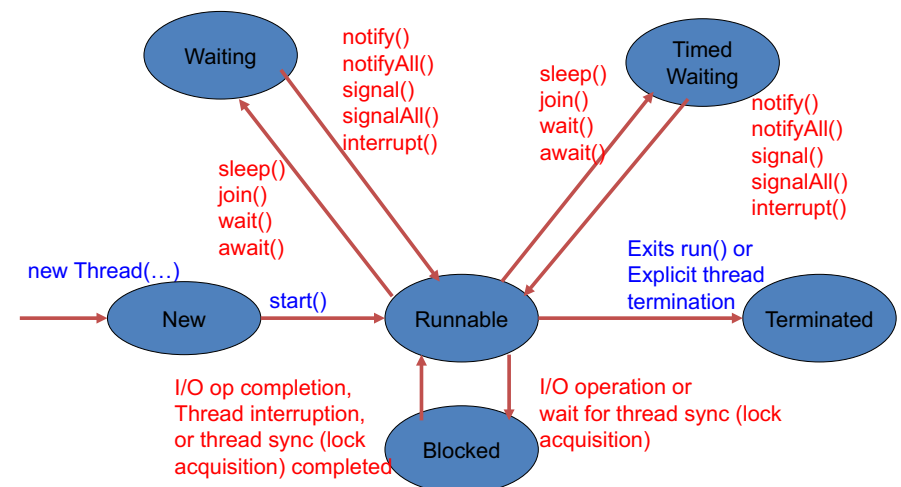
- Used to synchronize/serialize the threads that manipulate shared data.
- `java.util.concurrent.locks.Lock` interface
  - `ReentrantLock` class: the most commonly-used class for locking
    - Defines methods that allows threads to access shared data in a synchronized/serialized way.
- Atomic code is surrounded by `lock()` and `unlock()` method calls.

```
ReentrantLock aLock = new ReentrantLock();
aLock.lock();
atomic code
aLock.unlock();
```

22

# States of a Thread

- Once a thread calls `lock()`,
  - it acquires and owns a *lock* until it calls `unlock()`.
  - No other threads can acquire the lock until it is released by `unlock()`.
    - No other threads can execute atomic code until the lock is released with `unlock()`.
- If a thread calls `lock()` when another thread already owns the lock,
  - it goes to the *blocked* state and *gets blocked* (cannot do anything further) until the lock is released.



23

24

## How Can a Blocked Thread Run Again?

- JVM's thread scheduler
  - Periodically reactivates all blocked threads so that they can try to acquire the target lock.
    - If the lock is still unavailable, they get blocked again.
  - Detects a release of the target lock (i.e. completion of atomic code).
    - May notify all blocked threads so that one of them can acquire the target lock.
    - May choose one of the blocked threads to acquire the lock.
- Each blocked thread can eventually acquire the target lock when it is available.

25

## Coding Idiom for Locking

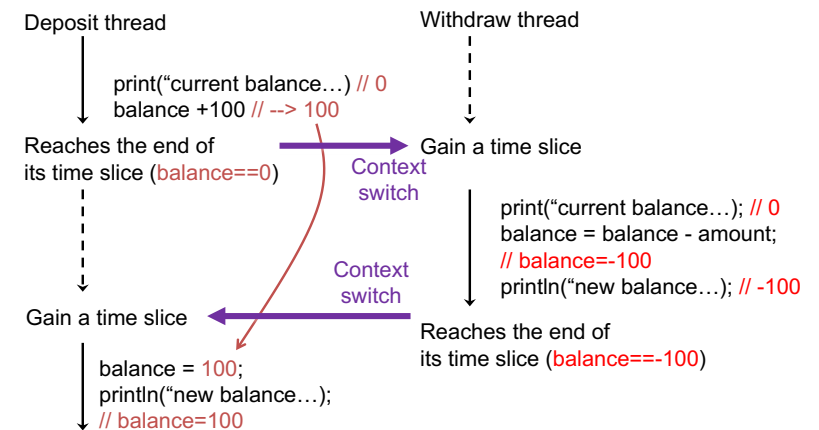
- Call `unlock()` in a finally clause.
  - `ReentrantLock aLock = new ReentrantLock();`  
`aLock.lock();`  
`try {`  
    *atomic code*  
`}`  
`finally {`  
    [aLock.unlock\(\);](#)  
`}`
- `unlock()` is never invoked
  - if `run()` returns from atomic code
  - if atomic code throws an exception
- A deadlock occurs.
  - Atomic code is locked forever, and no other threads can acquire the lock to execute the atomic code.

26

- |                                                                                                                                                                                                                                                              |                                                                                                                                                                                                                                                                           |
|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <ul style="list-style-type: none"><li><code>aLock.lock();</code><br/><code>try{</code><br/>    <i>atomic code</i><br/><code>}</code><br/><code>finally{</code><br/>    <a href="#">aLock.unlock();</a><br/><code>}</code><br/><br/><b>DO THIS!</b></li></ul> | <ul style="list-style-type: none"><li><code>try{</code><br/>    <a href="#">aLock.lock();</a><br/>    <i>atomic code</i><br/><code>}</code><br/><code>finally{</code><br/>    <a href="#">aLock.unlock();</a><br/><code>}</code><br/><br/><b>DON'T DO THIS!</b></li></ul> |
|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|

- Make sure to call `lock()` before a "try" block.
- If a thread throws an exception in `lock()`, it will not acquire the lock. However, it will call `unlock()`.
  - `lock()` can throw an `InterruptedException` when another thread call `interrupt()`.

## When Could a Context Switch Occur?



## When Could a Context Switch Occur?

- Across atomic operations.
- In a compound operation.

## Atomicity of Operations for Primitive Types

- The read and write operations for primitive data types, except double and long (64-bit) types, are atomic.
  - An “atomic” operation is transformed to a single bytecode instruction for a JVM.
  - While a thread works on an atomic operation, no other threads can work on it.
- `int x;`
  - Thread A does: `x=1`; Thread B does: `x=2`;
  - An assignment of an int value is atomic.
  - x contains 1 or 2 depending on which thread performs assignment earlier.
    - x never contains other values (e.g., 0 and 3) or corrupted data.
    - An example of corrupted data
      - » Some part of x (e.g., the first 16-bit of x) comes from Thread A and the remaining part (e.g., the other 16bit of x) comes from Thread B.

30

## Compound Operations

- A *compound* of atomic operations is NOT atomic.
  - `int i; boolean done;`
  - `done = true;` // 2 steps
  - `i = 1;` // 2 steps
  - `if(done)` // 2 steps
  - `i = j;` // 2 steps
  - `j = i + 1;` // 5 steps
    - Reading the value of i, reading/loading the value of 1, doing i+1, storing the result of i+1 to a certain memory space, and assigning the result to j.
  - `i = i + 1;` // 5 steps
  - `i++` // 5 steps
- A race condition can occur due to a context switch *in between* atomic operations/steps.

## An Example Race Condition

- `i = i + 1`
    - A compound of 5 atomic operations.
    - There are 4 places where race conditions can occur.
  - Thread synchronization enables **serialized** (or **atomic** or **exclusive**) access to a compound operation.
    - Allows **only one thread** to perform a compound operation at a time.
- ```
– ReentrantLock aLock = new ReentrantLock();
  aLock.lock();
  try{
    i = i + 1; // treated as an atomic operation
  }
  finally{
    aLock.unlock();
  }
```


Another Example

- ```
public void deposit(double amount){
 balance = balance + amount;
}
```

  - A compound of 5 atomic operations.
  - There are 4 places where race conditions can occur.
- Thread synchronization enables **serialized** (or **atomic** or **exclusive**) access to a compound operation.
  - Allows **only one thread** to perform a compound op at a time.
- ```
ReentrantLock aLock = new ReentrantLock();  
aLock.lock();  
try{  
    balance = balance + amount; // treated as an atomic op  
}  
finally{  
    aLock.unlock();  
}
```

Atomicity of Operations for Reference Types

- The read and write operations reference types are atomic.
 - A compound of atomic operations
 - e.g., `Foo foo = temp;` // requires multiple steps
 - A race condition can occur due to a context switch in between atomic steps.

What about 64-bit Types?

- The read and write operations for double and long variables are NOT atomic.
 - `long x;`
 - Thread A does: `x = 1L;`
 - Thread B does: `x = 2L;`
 - No guarantee that `x` contains 1L or 2L.
 - `x` can contain another value (e.g., 0L or 3L) or corrupted data.
 - `aLongVar = 100L;` // 2+ bytecode instructions
 - `if(aLongVar)` // 2+ bytecode instructions
 - `aLongVar ++` // 5+ bytecode instructions

34

ThreadSafeBankAccount

- Output
 - Lock obtained
 - Current balance (d): 0.0, New balance (d): 100.0
 - Lock released
 - Lock obtained
 - Current balance (w): 100.0, New balance (w): 0.0
 - Lock released
 - Lock obtained
 - Current balance (d): 0.0, New balance (d): 100.0
 - Lock released
 - Lock obtained
 - Current balance (w): 100.0, New balance (w): 0.0
 - Lock released

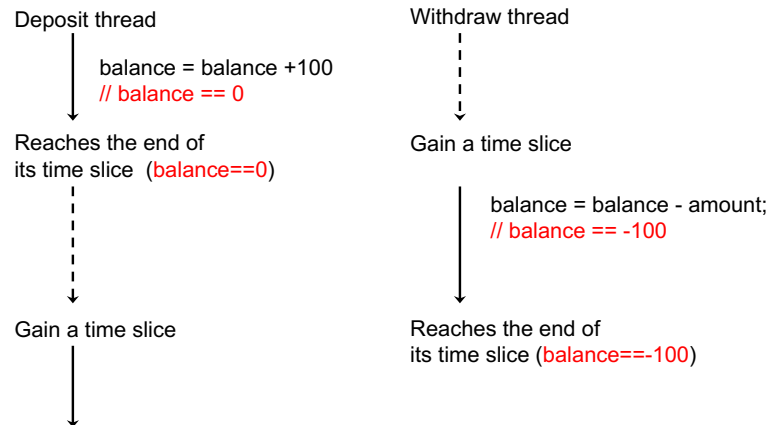
36

Be VERY Careful

- When multiple threads share and access a variable concurrently.
 - Make sure to guard the shared variable with a lock.
 - Surround reading and writing parts with lock() and unlock().
 - To treat the reading/writing parts as atomic code
- When a loop performs a conditional check with a shared variable (e.g., flag).
 - Surround reading part (i.e., conditional block) and writing part (i.e., flag-flipping part) with lock() and unlock()
 - To treat the reading/writing parts as atomic code

Consider this Lucky Case

- ```
public void deposit(double amount){
 balance = balance + amount; }
```



## What's Tricky in Thread Programming

- Your test code may or may not be able to detect race conditions.
  - It may not be able to detect race conditions even if you run it a lot of (e.g. few hundred) times.

## Nested Locking

- ```
class BankAccount {  
    private double balance;  
    private ReentrantLock lock;  
  
    public void deposit(double amount) {  
        lock.lock();  
        balance += amount; // 5 atomic steps  
        if (balance < MIN_BALANCE) // 4 atomic steps  
            subtractPenaltyFee();  
        lock.unlock(); }  
  
    private void subtractPenaltyFee() {  
        balance -= PENALTY; // 5 atomic steps  
        // NO NEED TO SURROUND THIS LINE BY LOCK() and UNLOCK()  
        // It is a part of atomic code.  
    }  
}
```

Thread Reentrancy

- ```
class A{
 private B b;
 private ReentrantLock lock;
 public void a1(){
 b = new B();
 lock.lock();
 b.b1(this); //nested locking
 lock.unlock();
 }
 public void a2(){
 lock.lock();
 do something.
 lock.unlock();
 }
}
```
  - ```
Class B{  
    public void b1(A a){  
        a.a2();  
    }  
}
```
- If a thread performs:
`A a = new A();`
`a.a1();`
it *re-enters (or re-acquires)* the same lock that it already owns.
- This code does not have a deadlock problem.
 - A thread can *re-enter* the same lock as far as it already owns the lock.