

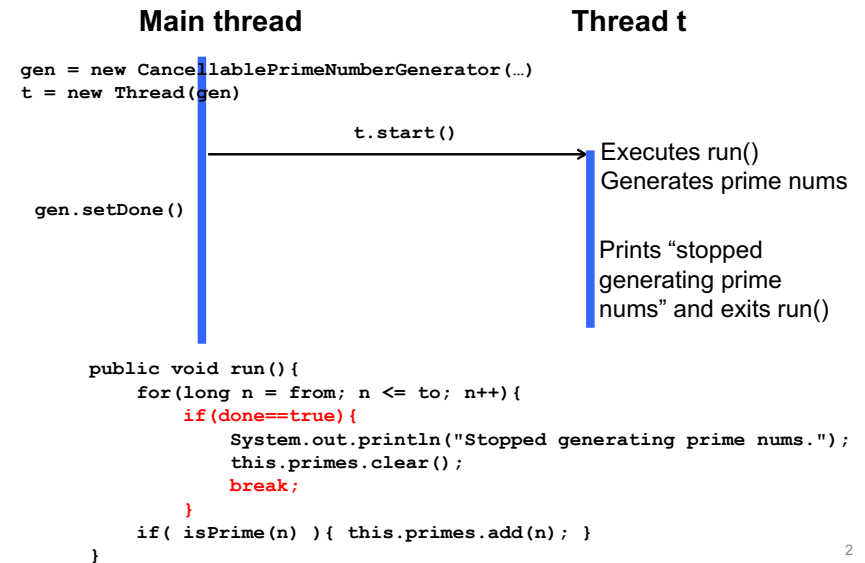
CancellablePrimeNumberGenerator

```
class CancellablePrimeNumberGenerator extends PrimeNumberGenerator{
    private boolean done = false;

    public void run(){
        for( long n = from; n <= to; n++ ){
            if(done==true){
                System.out.println("Stopped generating prime nums.");
                this.primes.clear();
                break;
            }
            if( isPrime(n) ){ this.primes.add(n); }
        }
    }

    public void setDone(){
        done = true;
    }
}
```

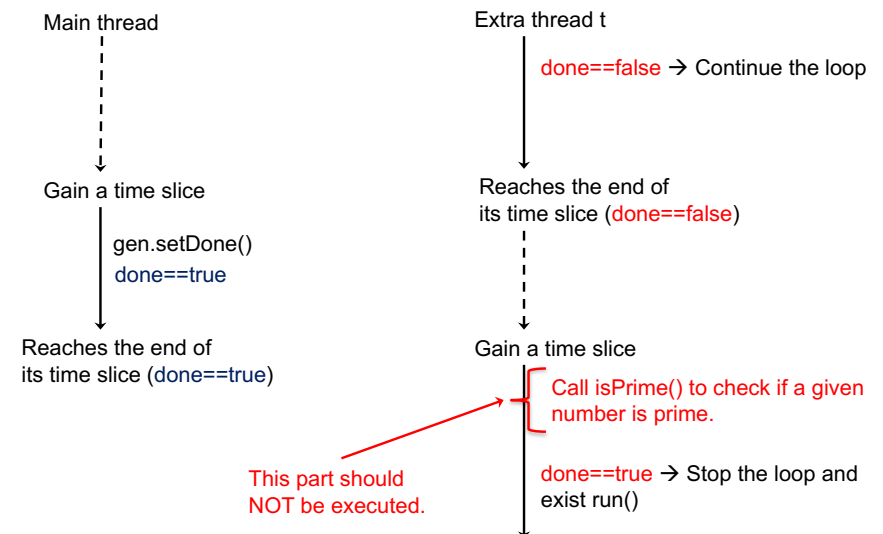
1



2

- Not thread-safe. Race conditions can occur.
 - Thread safety:
 - No race conditions
 - No deadlocks

A Potential Race Condition



4

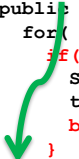
```

class CancellablePrimeNumberGenerator extends PrimeNumberGenerator{
    private boolean done = false;

    public void run(){
        for( long n = from; n <= to; n++ ){
            if(done==true){
                System.out.println("Stopped generating prime nums.");
                this.primes.clear();
                break;
            }
            if( isPrime(n) ){ this.primes.add(n); }
        }

        public void setDone(){
            done = true;
        }
    }
}

```



5


```

class CancellablePrimeNumberGenerator extends PrimeNumberGenerator{
    private boolean done = false;

    public void run(){
        for( long n = from; n <= to; n++ ){
            if(done==true){
                System.out.println("Stopped generating prime nums.");
                this.primes.clear();
                break;
            }
            if( isPrime(n) ){ this.primes.add(n); }
        }

        public void setDone(){
            done = true;
        }
    }
}

```



6

Visibility Issue

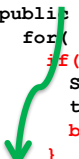

```

class CancellablePrimeNumberGenerator extends PrimeNumberGenerator{
    private boolean done = false;

    public void run(){
        for( long n = from; n <= to; n++ ){
            if(done==true){
                System.out.println("Stopped generating prime nums.");
                this.primes.clear();
                break;
            }
            if( isPrime(n) ){ this.primes.add(n); }
        }

        public void setDone(){
            done = true;
        }
    }
}

```

7

- The current (most up-to-date) value of the shared variable “done” is not visible for all threads.
- **ReentrantLock** allows your code to preserve atomicity AND visibility.

Solution: Locking and Balking

- A General form of the “balking” idiom

```
- boolean done = false;
  ReentrantLock lock = new ReentrantLock();
  ...
  while(true){
    lock.lock();
    try{
      if(done) break; // Balk
      ...           // Do some task
    }finally{
      lock.unlock();
    }
    ...
  }

- void setDone(){
  lock.lock();
  try{
    done = true;
  }finally{
    lock.unlock();
  }
}
```

- Threads must use the same instance of `ReentrantLock`.

- Try NOT to surround the entire loop with `lock()` and `unlock()`!
- Why?
 - May result in a deadlock.
 - Does not enjoy concurrency.

Be Careful for Potential Race Conditions

- When multiple threads share and access a variable concurrently.
 - Make sure that a shared variable is guarded with a lock.
 - Surround reading and writing parts with `lock()` and `unlock()`.
 - Reading/writing parts → atomic code
 - c.f. Thread-unsafe and -safe bank accounts
- When a loop performs a conditional check with a shared variable (i.e., flag).
 - Surround reading part (i.e., conditional block) and writing part (i.e., flag-flipping part) with `lock()` and `unlock()`
 - Reading/writing parts → atomic code
 - Try NOT to surround the entire loop with `lock()` and `unlock()`!

Treating the Entire Loop as Atomic Code May Result in a Deadlock

- | | |
|---|---|
| <ul style="list-style-type: none">• DO NOT do this.• <pre>try{ lock.lock(); while(!done){ // Do some task System.out.println("#"); } }finally{ lock.unlock(); }</pre>• <pre>lock.lock(); try{ done = true; }finally{ lock.unlock(); }</pre> | <ul style="list-style-type: none">• Do this.• <pre>while(true){ lock.lock(); try{ if(done) break; // Balk // Do some task System.out.println("#"); }finally{ lock.unlock(); } }</pre>• <pre>lock.lock(); try{ done = true; }finally{ lock.unlock(); }</pre> |
|---|---|

Treating the Entire Loop as Atomic Code May Result in a Deadlock

- If a thread acquires the lock and starts printing #s, it will print #s forever.
 - No other threads cannot flip the flag forever (deadlock!)

```
- lock.lock();
while(!done){
    System.out.println("#"); // reading part
}
lock.unlock();
```

The purple thread gets stuck here forever because the green thread never release the lock.

```
- lock.lock();
done = true;
lock.unlock(); // writing part
```

Treating the Entire Loop as Atomic Code Does NOT Enjoy Concurrency

- DO NOT do this.


```
long n;
try{
    lock.lock();
    for(n = from; n <= to; n++){
        if(done==true) break;
        if(isPrime(n)){
            this.primes.add(n);
        }
    }
}finally{
    lock.unlock();
}
```
- Do this.


```
long n;
for(n = from; n <= to; n++){
    try{
        lock.lock();
        if(done==true) break;
        if(isPrime(n)){
            this.primes.add(n);
        }
    }finally{
        lock.unlock();
    }
}
```
- lock.lock();


```
try{
    done = true;
}finally{
    lock.unlock();
}
```
- lock.lock();


```
try{
    done = true;
}finally{
    lock.unlock();
}
```

Treating the Entire Loop as Atomic Code Does NOT Enjoy Concurrency

- If a thread acquires the lock and starts generating prime numbers, it will release the lock when $n > to$.
 - No other threads cannot flip the flag until the loop ends.
 - No deadlock occurs because the loop ends when $n > to$.

```
- try{
    lock.lock();
    for(n = from; n <= to; n++){
        if(done==true) break;
        if(isPrime(n)){
            this.primes.add(n);
        }
    }
}finally{
    lock.unlock();
}
```

The purple thread can acquire the lock after all prime numbers have been generated.

```
- lock.lock();
try{
    done = true;
}finally{
    lock.unlock();
}
```

Treating the Entire Loop as Atomic Code Does NOT Enjoy Concurrency

- This code is thread-safe, but it does not enjoy concurrency.
 - While the green thread generates prime numbers for a given range in between "from" and "to," the purple thread cannot stop the green thread.


```
- try{
    lock.lock();
    for(n = from; n <= to; n++){
        if(done==true) break;
        if(isPrime(n)){
            this.primes.add(n);
        }
    }
}finally{
    lock.unlock();
}
```

lock() returns when the green thread releases the lock.

```
- lock.lock();
try{
    done = true;
}finally{
    lock.unlock();
}
```



HW 7

- Revise CancellablePrimeNumberGenerator.java to make it thread-safe.
 - Use ReentrantLock
 - Use try-finally blocks.
 - Call unlock() in a finally block. Always do this in all subsequent HWs.
 - Use balking
 - Do not surround the “for” loop with lock() and unlock().



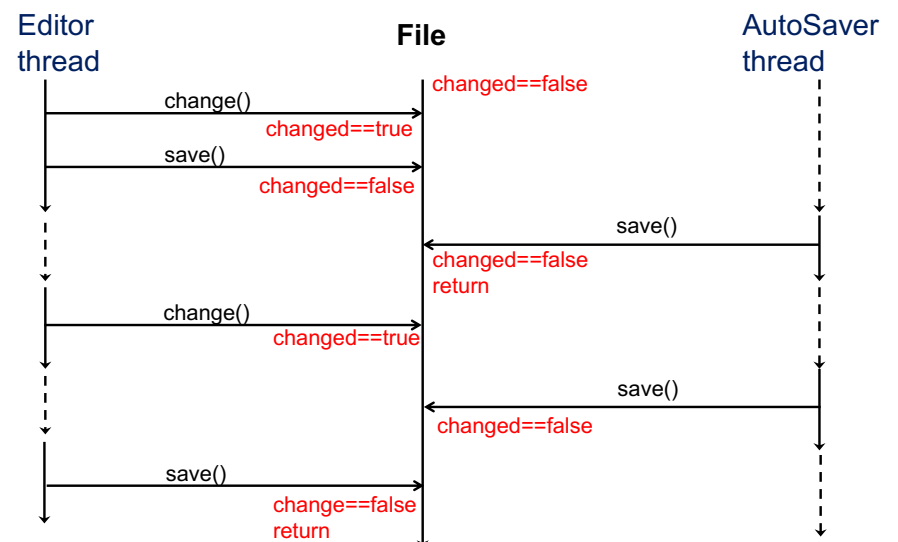
- File
 - Has a boolean variable: “changed”
 - Initialized to be false.
 - change()
 - Changes the file’s content.
 - Assigns true to the variable “changed.”
 - save()
 - if(changed==false) return;
 - if(changed==true)
 - print out some message (e.g., time stamp, etc.)
 - assigns false to the variable “changed.”
- Editor (a Runnable) repeats:
 - Calls change() and save()
 - Sleeps for a second.
- AutoSaver (a Runnable) repeats:
 - Calls save()
 - Sleeps for two seconds.

Exercise: Concurrent Access to a File



- Imagine word processing software.
- Assume two threads
 - One for a file editing feature
 - Allows the user to edit a file and save it.
 - One for an automatic file saving feature
 - Periodically saves an open file at background.
- The 2 threads call `change()` and `save()` on an open file concurrently.

Desirable Result



HW 8

- Race conditions can occur if you do not guard the “changed” variable with a lock. Explain a potential race condition with a diagram like in the previous slide.
- Submit thread-safe code.
 - Define 3 classes: `File`, `Editor` and `AutoSaver`
 - Define a lock in `File`. Use the lock in `change()` and `save()`
 - Use try-finally blocks: Always do this in all subsequent HWs.
 - Create two extra threads and have them execute `Editor`’s `run()` and `AutoSaver`’s `run()`
 - Those threads acquire and release the lock via `change()` and `save()`.

- Have the main thread sleep for some time while `Editor` and `AutoSaver` are running.
 - Use `Thread.sleep()`
- Have the main thread terminate the other two threads.
 - Define a flag variable “done” and `setDone()` in `Editor` and `AutoSaver`
- Note that this code is not thread-safe.
 - c.f. `CancellablePrimeNumberGenerator.java`
- Use a lock in `Editor` and `AutoSaver` to guard flag variables.
 - Use try-finally blocks
 - Use balking in `run()`
 - Do not surround “while” loops with `lock()` and `unlock()`.

```

class Editor implements Runnable{
    private boolean done = false;

    public void run(){
        while(true){
            if(done==true){
                System.out.println("...");
                break;
            }
            aFile.change();
            aFile.save();
            Thread.sleep(1000);
        }
    }

    public void setDone(){
        done = true;
    }
}

```

Recap: Singleton Design Pattern

- Guarantee that a class has only one instance.

```

public class Singleton{
    private Singleton(){};
    private static Singleton instance = null;

    // Factory method to return a singleton instance
    public static Singleton getInstance(){
        if(instance==null)
            instance = new Singleton();
        return instance;
    }
}

```

HW 9

- The Singleton class is not thread-safe.
 - Race conditions can occur if you do not guard the “instance” variable with a lock. Explain a potential race condition where more than one instances are created.
 - Use a diagram like in a previous slide.
- Submit a thread-safe version of Singleton.
 - Define a lock in Singleton. Use the lock in `getInstance()`
 - Use try-finally blocks: Always do this in all subsequent HWs.
 - Create multiple extra threads and have them call `getInstance()`
 - Make sure that only one instance is created.
 - Use `System.out.println(Singleton.getInstance())`

Concurrent Singleton Design Pattern

- Guarantee that a class has only one instance.
- ```
public class Singleton{
 private Singleton(){};
 private static Singleton instance = null;
 private static ReentrantLock lock = new ReentrantLock();

 // Factory method to create or return a singleton instance
 public static Singleton getInstance(){
 lock.lock();
 if(instance==null)
 instance = new Singleton();
 lock.unlock();
 return instance;
 }
}
```

25

## Exercise: Regular and Static Locks

- ```
public class Foo{  
    private ReentrantLock lock = new ReentrantLock();  
    private static ReentrantLock sLock = new ReentrantLock();  
  
    public void a() {...}  
    public void b() {...}  
    public void syncA() {lock.lock(); ... lock.unlock();}  
    public void syncB() {lock.lock(); ... lock.unlock();}  
  
    public static void sA() {...}  
    public static void sB() {...}  
    public static void sSyncA() {sLock.lock(); ... sLock.unlock();}  
    public static void sSyncB() {sLock.lock(); ... sLock.unlock(); }  
}
```
- `x = new Foo(); y = new Foo();`
- Two threads call...
 - x.a() and x.a(): no synchronization (no exclusive execution) for the two threads
 - x.a() and x.b(): no synchronization
 - x.a() and x.syncA(): no synchronization
 - x.syncA() and x.syncA(): Synchronization (exclusive execution)
 - y.syncA() and y.syncB(): Synchronization
 - x.syncA() and y.syncA(): No synchronization
 - x.syncA() and y.syncB(): No synchronization

Regular and Static Locks


- ```
public class Foo{
 ReentrantLock lock = new ReentrantLock();
 static ReentrantLock sLock = new ReentrantLock(); }
```
- A **regular lock** is created and used on an instance-by-instance basis.
  - Different instances of Foo have different locks (i.e. different instances of ReentrantLock).
- A **static lock** is created and used on a per-class basis.
  - All instances of Foo share a single lock (“sLock”).

- ```
public class Foo{  
    private ReentrantLock lock = new ReentrantLock();  
    private static ReentrantLock sLock = new ReentrantLock();  
  
    public void a() {...}  
    public void b() {...}  
    public void syncA() {lock.lock(); ... lock.unlock();}  
    public void syncB() {lock.lock(); ... lock.unlock();}  
  
    public static void sA() {...}  
    public static void sB() {...}  
    public static void sSyncA() {sLock.lock(); ... sLock.unlock();}  
    public static void sSyncB() {sLock.lock(); ... sLock.unlock(); }  
}
```
- `x = new Foo(); y = new Foo();`
- Two threads call...
 - x.a() and Foo.sA(): No synchronization for the two threads
 - x.syncA() and Foo.sA(): No synchronization
 - Foo.sA() and Foo.sA(): No synchronization
 - Foo.sA() and Foo.sB(): No synchronization
 - x.syncA() and Foo.sSyncA(): No synchronization
 - Foo.sSyncA() and Foo.sSyncA(): Synchronization
 - Foo.sSyncA() and Foo.sSyncB(): Synchronization
 - x.sSyncA() and y.sSyncB(): Synchronization
 - This is not grammatically wrong, but write Foo.sSyncA() instead of x.sSyncA()

Thread.sleep()

- ```
Thread t = new Thread(new FooRunnable());
t.start();
try{
 t.sleep(1000);
}catch(InterruptedException e){...}
```
- It looks like an extra thread (*t*) will sleep.
- However, the main thread will actually sleep
  - because `sleep()` is a **static method** of `Thread`.
  - `Thread.sleep()`: Causes the currently executing thread to sleep (temporarily cease execution) for the specified number of milliseconds
- DO NOT write `t.sleep(...)`. It's misleading and error-prone.
- ALWAYS WRITE `Thread.sleep(...)`.
  - Make sure to do this in HW 8.

## Another Exercise



- ```
class SecurityGate{
    private int counter = 0;

    public void enter(){
        counter++;
    }

    public void exit(){
        counter--;
    }

    //Get the # of guests
    //in the gate
    public int getCount(){
        return counter;
    }
}
```
- ```
class Guest implements Runnable{
 private SecurityGate gate;

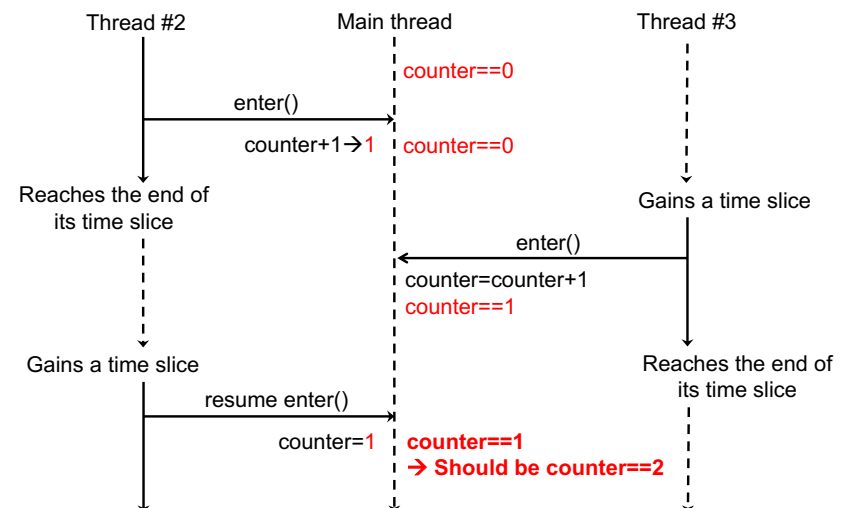
 public Guest(.....){
 gate = SecurityGate.getInstance();
 }

 public void run(){
 gate.enter();
 gate.getCount();
 gate.exit();
 }
}
```


## Not Thread Safe!

- `SecurityGate` is not thread-safe due to potential race conditions.
- `counter++` is a compound operation and not atomic.
  - Syntactic sugar for `counter = counter + 1;`
  - Requires multiple atomic operations
    - A context switch can occur across different atomic ops.
- The same goes to `counter--`.

## A Potential Race Condition







## HW 10-1

```

• class SecurityGate{
 private int counter = 0;
 private ReentrantLock lock;

 public void enter(){
 lock.lock();
 counter++;
 lock.unlock();
 }

 public void exit(){
 lock.lock();
 counter--;
 lock.unlock();
 }

 public int getCount(){
 return counter;
 }
}

• class Guest implements Runnable{
 private SecurityGate gate;

 public Guest(.....){
 gate = SecurityGate.getInstance();
 }

 public void run(){
 gate.enter();
 gate.exit();
 gate.getCount();
 }
}

```

- Make `SecurityGate` thread-safe.
  - Revise `SecurityGate` to be a **concurrent singleton class**.
  - Use two locks in `SecurityGate`
    - One to guard counter
      - Use it in `enter()` and `exit()`
    - One to guard instance
      - Use it in `getInstance()`
- Create many threads to execute `Guest` instances' `run()` and have them call `enter()`, `exit()` and `getCount()`.

## An Alternative Solution: Use `AtomicInteger`

- Offers a series of thread-safe methods to manipulate an integer value **atomically**.
  - `int i;`  
`i = i + 1;` // Not thread-safe
  - `AtomicInteger atomicInt = new AtomicInteger(0);`  
`atomicInt.incrementAndGet();` // Thread-safe
  - `decrementAndGet(), addAndGet(int), set(), get(), getAndSet(int), compareAndSet(int, int)...`
  - `updateAndGet(IntUnaryOperator), accumulateAndGet(int, IntBinaryOperator)`
- Many of the methods do not use locking; they are faster than lock-based code.

## Be Careful!

- ```

• int i = ...
  i = i + 1; // Not thread-safe

• AtomicInteger atomicInt = new AtomicInteger(0);
  atomicInt.incrementAndGet(); // Thread-safe
  – This is atomic.

• AtomicInteger atomicInt = new AtomicInteger(0);
  int i = atomicInt.incrementAndGet(); // NOT thread-safe
  – This is compound!
    • incrementAndGet() is atomic though.
  
```

AtomicInteger

– `updateAndGet(IntUnaryOperator updateFunction)`

- `IntUnaryOperator`: a general-purpose functional interface
- Atomically updates the current integer value with the result of applying the given function, returning the updated value.
- `updateFunction`: a lambda expression that takes the current int value as a parameter, updates it and returns the updated value.
- ```
AtomicInteger atomicInt = new AtomicInteger(10);
atomicInt.updateAndGet((int i)->++i); // 11. Thread safe
atomicInt.incrementAndGet(); // 12. Thread safe
```

|                                     | Params         | Returns        | Example use case |
|-------------------------------------|----------------|----------------|------------------|
| <code>UnaryOperator&lt;T&gt;</code> | <code>T</code> | <code>T</code> | Logical NOT (!)  |

## – `accumulateAndGet(int, IntBinaryOperator)`

- `IntBinaryOperator`: a general-purpose functional interface
- Atomically updates the current integer value with the result of applying the given function, returning the updated value.
- ```
atomicInt.accumulateAndGet(initValue,
    (result, currentVal)-> ... );
```
- ```
int result = initValue;
result = accumulate(result, currentVal);
```
- Takes a lambda expression as the second parameter.
  - The body of `accumulate()` is expressed in the LE.
  - C.f. `Stream.reduce()`

|                                      | Params            | Returns        | Example use case            |
|--------------------------------------|-------------------|----------------|-----------------------------|
| <code>BinaryOperator&lt;T&gt;</code> | <code>T, T</code> | <code>T</code> | Multiplying two numbers (*) |

- ```
AtomicInteger atomicInt = new AtomicInteger(10);
atomicInt.updateAndGet( (int i)->++i ); // 11. Thread safe
```

• Why ++1? Just in case, note that:

```
- int i = 0;
  i++; // i==1

- int i = 0;
  int x = i++; // i==1, x==0

- int i = 0;
  int y = ++i; // i==1, y==1
```

- ```
AtomicInteger atomicInt = new AtomicInteger(0);
atomicInt.accumulateAndGet(10,
 (result, currentVal)->
 (currentVal+result)/2);
```
- ```
AtomicInteger atomicInt = new AtomicInteger(0);
atomicInt.accumulateAndGet(
    10, (result, currentVal)->{
        if(currentVal >= result) return currentVal;
        else if return result; }));
```
- ```
int currentVal = 0; int result = 10;
result = accumulate(result, currentVal);
```

## Again...

- ```
AtomicInteger atomicInt = new AtomicInteger(0);
atomicInt.accumulateAndGet(10,
    (result, currentVal)->
        (currentVal+result)/2 );
```

 - This is atomic.
- ```
AtomicInteger atomicInt = new AtomicInteger(0);
int average = atomicInt.accumulateAndGet(
 10, (result, currentVal)->(currentVal+result)/2);
```

  - This is compound!
    - The methods of `AtomicInteger` are atomic though.

## HW 10-2

- Define the “counter” variable with `AtomicInteger`.
- Increment and decrement the “counter” variable with `AtomicInteger.updateAndGet()`
  - Do NOT use `incrementAndGet()` and `decrementAndGet()`.
- Create many threads to execute Guest instances’ `run()` and have them call `enter()`, `exit()` and `getCount()`.

## java.util.concurrent.atomic Package

- `AtomicBoolean`
- `AtomicInteger`
- `AtomicIntegerArray`
- `AtomicLong`
- `AtomicLongArray`
- `AtomicReference<V>`
- `AtomicReferenceArray<E>`
- `DoubleAccumulator`
- `DoubleAdder`
- `LongAccumulator`
- `LongAdder`
- ...

## Where did the synchronized Keyword go?

- Java still has the `synchronized` keyword.
  - ```
public synchronized void foo(){
    // The entire method body is atomic.
}
```
 - ```
public void foo(){
 // non-atomic code here
 synchronized(this){
 // atomic code here
 }
 // non-atomic code here }
```
- Implicit locking
  - No need to create a `ReentrantLock` and call `lock()` and `unlock()`.
- When a thread enters a synchronized method/block, it tries to acquire the (implicit) lock that `this` instance maintains.
  - Instance-by-instance locking
- Code gets tricky/dirty to use multiple locks in a single class.

- Explicit locking

- ```
ReentrantLock aLock = new ReentrantLock();  
public void foo(){  
    aLock.lock();  
    // atomic code  
    aLock.unlock(); }
```

- Arbitrary locking scope.
 - Clean code even if a class uses multiple locks.
 - Extra functionalities
 - e.g., `getQueueLength()`: returns the # of waiting threads.
 - `tryLock()`: acquires a lock only if it is not held by another thread.
 - The catch is... it's VERY easy to forget calling `unlock()`.
 - Must call `unlock()` in a finally clause.

- Implicit locking with the “synchronized” keyword
 - A thread can call `notify()` and `notifyAll()` even if it has not acquired a lock.
 - An `IllegalMonitorStateException` is thrown.

- Explicit locking

- This error/bug never occurs.

- ```
ReentrantLock lock = new ReentrantLock();
Condition cond = lock.newCondition();
lock.lock();
...
cond.signalAll();
```