


Thread Safety Issues

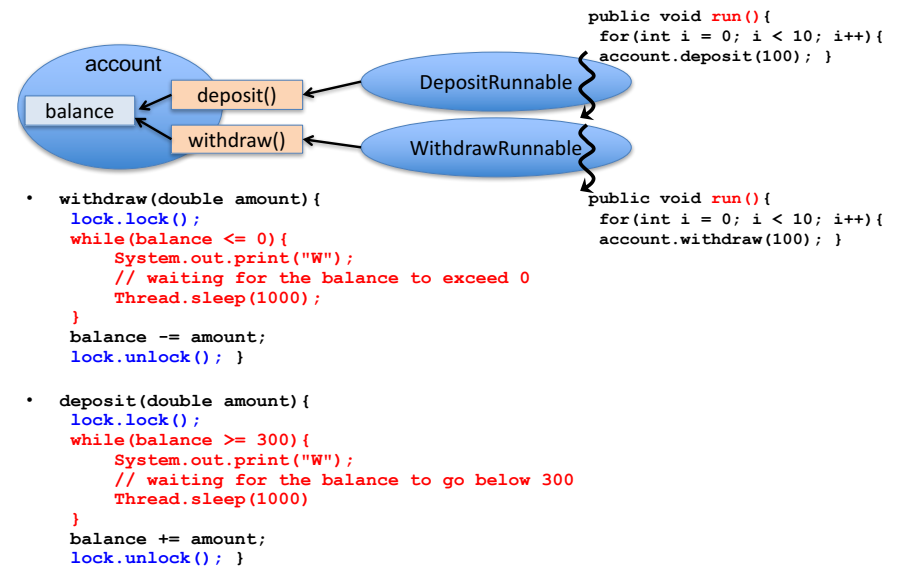
- Race conditions
- Deadlocks
- Thread-safe code is free from race conditions and deadlocks.

Deadlock

DeadlockedBankAccount.java

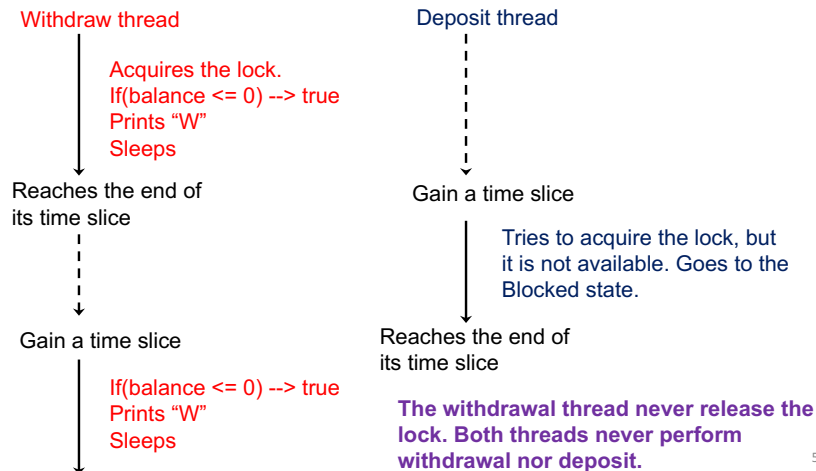


DeadlockedBankAccount.java



How Can a Deadlock Occur?

- Suppose the withdrawal thread goes ahead.



5

Note

- A JVM can perform context switches even when a thread runs atomic code.
 - A lock guarantees that only one thread exclusively runs atomic code at a time.
- Some resources explicitly/implicitly say that context switches never occur when a thread runs atomic code.
 - It is WRONG!

DeadlockedBankAccount2.java

- Previous version

```
- withdraw(double amount){
    lock.lock();
    while(balance <= 0){
        System.out.print("W");
        // waiting for the
        // balance to exceed 0
        Thread.sleep(1000);
    }
    balance -= amount;
    lock.unlock();
}

- deposit(double amount){
    lock.lock();
    while(balance > 300){
        System.out.print("W");
        // waiting for the balance
        // to go below 300
        Thread.sleep(1000);
    }
    balance += amount;
    lock.unlock(); }
```

- New version

```
- withdraw(double amount){
    while( balance <= 0 ){
        System.out.print("W");
        Thread.sleep(2);
    }
    lock.lock();
    balance -= amount;
    lock.unlock();
}

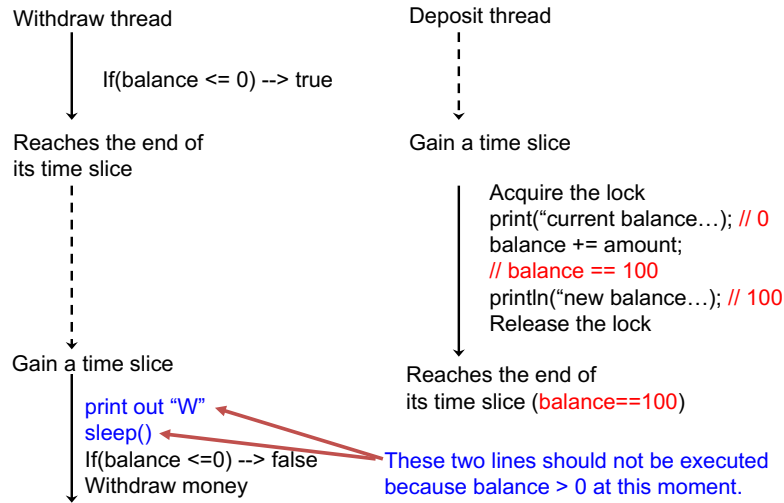
- deposit(double amount){
    while( balance > 300 ){
        System.out.print("W");
        Thread.sleep(2);
    }
    lock.lock();
    balance += amount;
    lock.unlock();
}
```

7

- Has no deadlock problems.
- Can generate race conditions.

8

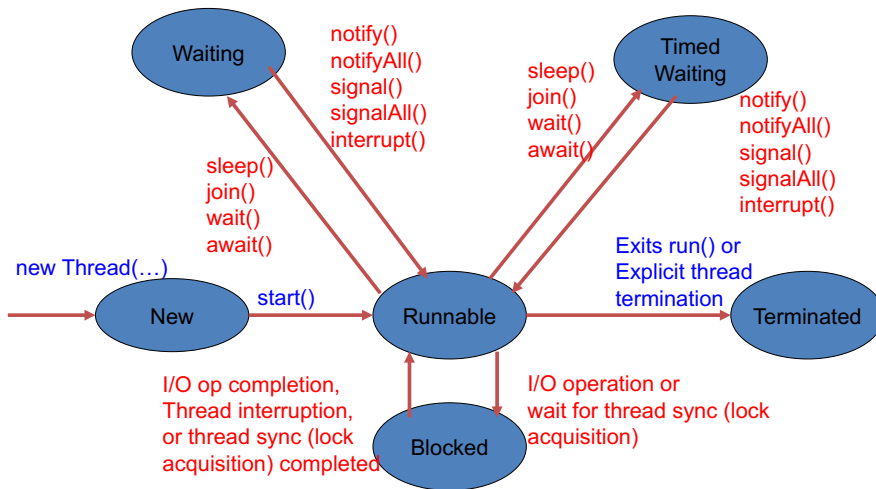
A Potential Race Condition in DeadlockedBankAccount2



Avoiding Deadlocks and Race Conditions

- Use a **Condition** object(s).
 - Allows a thread to
 - Temporarily release a lock so that another thread can acquire it and proceed.
 - Re-acquire the lock later.
- `java.util.concurrent.locks.Condition`
 - Obtain its instance from a lock object
 - ```
ReentrantLock lock = new ReentrantLock();
Condition condition = lock.newCondition(); // Factory method
condition.await(); // Temporarily releases the lock
// Goes to the Waiting state
```

## States of a Thread



## ThreadSafeBankAccount2.java

- Condition `sufficientFundsCondition = lock.newCondition();`  
Condition `belowUpperLimitFundsCondition = lock.newCondition();`
- `withdraw(double amount){`  
    `lock.lock();`  
    `while(balance <= 0){`  
        `// Wait for the balance to exceed 0`  
        `sufficientFundsCondition.await(); }`  
    `balance -= amount;`  
    `belowUpperLimitFundsCondition.signalAll();`  
    `lock.unlock(); }`
- `deposit(double amount){`  
    `lock.lock();`  
    `while(balance >= 300){`  
        `// Wait for the balance to go below 300.`  
        `belowUpperLimitFundsCondition.await(); }`  
    `balance += amount;`  
    `sufficientFundsCondition.signalAll();`  
    `lock.unlock(); }`

## ThreadSafeBankAccount2.java

```
• Condition sufficientFundsCondition = lock.newCondition();
 Condition belowUpperLimitFundsCondition = lock.newCondition();

• withdraw(double amount){
 lock.lock();
 while(balance <= 0){
 // Wait for the balance to exceed 0
 sufficientFundsCondition.await();
 }
 balance -= amount;
 belowUpperLimitFundsCondition.signalAll();
 lock.unlock();
}

• deposit(double amount){
 lock.lock();
 while(balance >= 300){
 // Wait for the balance to go below 300.
 belowUpperLimitFundsCondition.await();
 }
 balance += amount;
 sufficientFundsCondition.signalAll();
 lock.unlock();
}
```

A "deposit" thread calls signalAll() to wake up a thread(s) that is/are waiting until balance > 0.

13

## ThreadSafeBankAccount2.java

```
• Condition sufficientFundsCondition = lock.newCondition();
 Condition belowUpperLimitFundsCondition = lock.newCondition();

• withdraw(double amount){
 lock.lock();
 while(balance <= 0){
 // Wait for the balance to exceed 0
 sufficientFundsCondition.await();
 }
 balance -= amount;
 belowUpperLimitFundsCondition.signalAll();
 lock.unlock();
}

• deposit(double amount){
 lock.lock();
 while(balance >= 300){
 // Wait for the balance to go below 300.
 belowUpperLimitFundsCondition.await();
 }
 balance += amount;
 sufficientFundsCondition.signalAll();
 lock.unlock();
}
```

A "withdraw" thread calls signalAll() to wake up a thread(s) that is/are waiting until balance < 300.

14

## ThreadSafeBankAccount2.java

```
• Condition sufficientFundsCondition = lock.newCondition();
 Condition belowUpperLimitFundsCondition = lock.newCondition();

• withdraw(double amount){
 lock.lock();
 while(balance <= 0){
 // Wait for the balance to exceed 0
 sufficientFundsCondition.await();
 }
 balance -= amount;
 belowUpperLimitFundsCondition.signalAll();
 lock.unlock();
}

• deposit(double amount){
 lock.lock();
 while(balance >= 300){
 // Wait for the balance to go below 300.
 belowUpperLimitFundsCondition.await();
 }
 balance += amount;
 sufficientFundsCondition.signalAll();
 lock.unlock();
}
```


15

## Condition

- **await()**
  - Will wait until it is signaled or interrupted
  - Will wait until it is signaled or interrupted, or until a specified waiting time (relative time) elapsed.
  - Will wait until it is signaled or interrupted, or until a specified deadline (absolute time).
  - If signaled, goes to the Runnable state and re-acquires a lock.
    - Will be "blocked" if the thread re-acquisition fails.
  - Throws an **InterruptedException** if interrupted.
    - c.f. A previous lecture note on thread interruption
- **signalAll()**
  - Wakes up all waiting threads on a condition object.
    - All of them go to the "runnable" state.
    - One of them will re-acquire a lock.

16

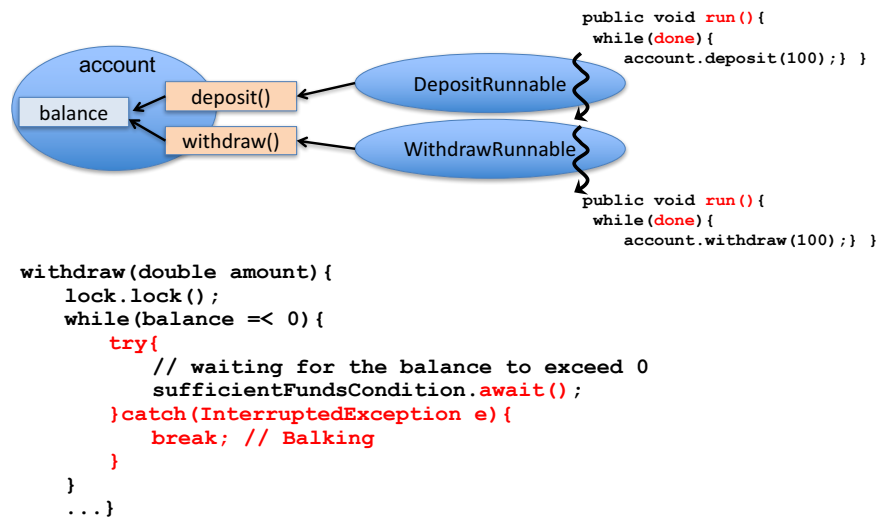
## States of a Thread



17

- When a thread calls `await()`, `signal()` or `signalAll()` on a Condition object,
  - the thread is assumed to hold a lock associated with the Condition object.
  - If the thread does not, an `IllegalMonitorStateException` is thrown.

## 2-Step Thread Termination (c.f. Lec Note #9)



## HW 13

- Implement 2-step termination for “deposit” and “withdraw” threads.
  - Implement a flag-based termination scheme in `DepositRunnable` and `WithdrawRunnable`
    - To let “deposit” and “withdraw” threads to return `run()`.
  - Have the main thread call `interrupt()` on “deposit” and “withdraw” threads
    - To let those threads to wake up in case they are in the Waiting state due to `await()` or `sleep()`.
- Due: April 24 (Tue) midnight

## signalAll() Before or After a State Change?

- ```

withdraw(double amount){
    lock.lock();
    while(balance <= 0){
        // waiting for the balance to exceed 0
        sufficientFundsCondition.await(); }
        balance -= amount;
        belowUpperLimitFundsCondition.signalAll();
    lock.unlock(); }

```
- ```

deposit(double amount){
 lock.lock();
 while(balance >= 300){
 // waiting for the balance to go below 300.
 belowUpperLimitFundsCondition.await(); }
 balance += amount;
 sufficientFundsCondition.signalAll();
 lock.unlock(); }

```
- What if you call signalAll() first and then update the balance? Will any thread safety issues come out?

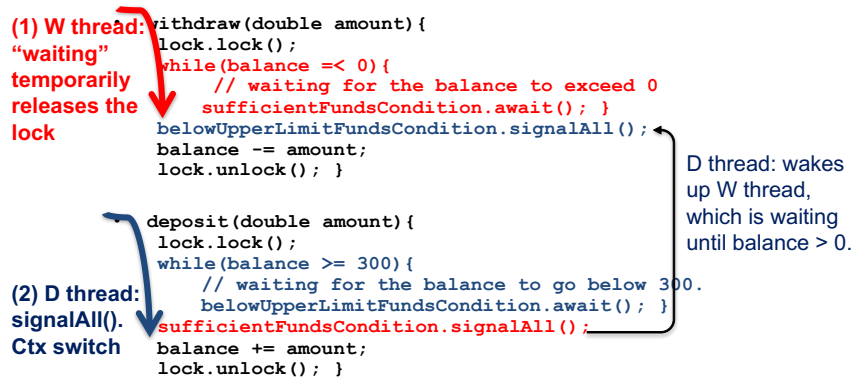
- ```

withdraw(double amount){
    lock.lock();
    while(balance <= 0){
        // waiting for the balance to exceed 0
        sufficientFundsCondition.await(); }
        belowUpperLimitFundsCondition.signalAll();
        balance -= amount;
    lock.unlock(); }

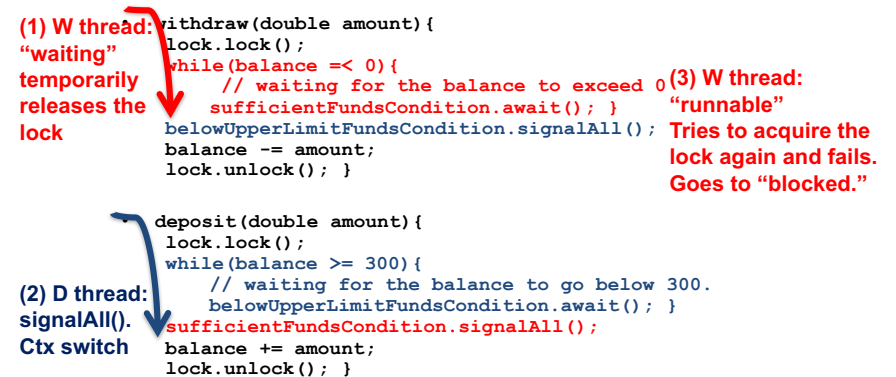
```
- ```

deposit(double amount){
 lock.lock();
 while(balance >= 300){
 // waiting for the balance to go below 300.
 belowUpperLimitFundsCondition.await(); }
 sufficientFundsCondition.signalAll();
 balance += amount;
 lock.unlock(); }

```
- For example, do you need to worry about race conditions in this case?

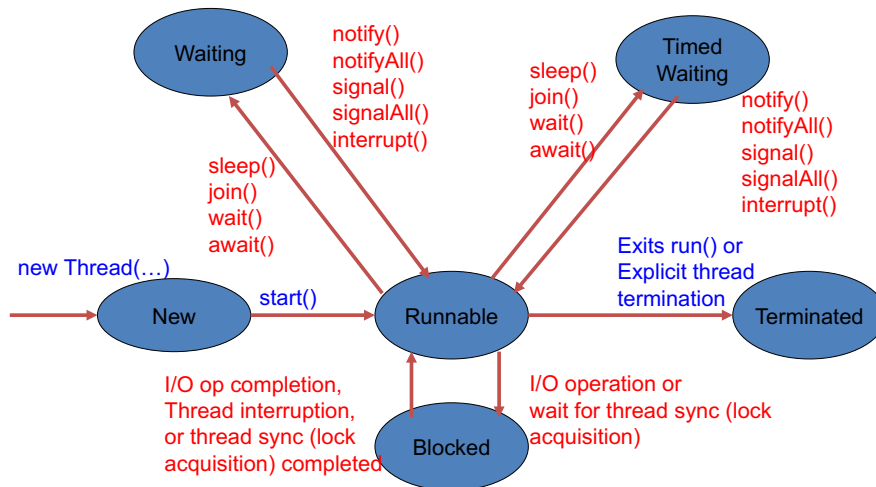


- Can the "W" thread withdraw money before the "D" thread deposits money?
  - Can the balance have a negative value?
    - The answer is NO.



- "W" thread CANNOT withdraw money before "D" thread deposits money.
- "D" thread CANNOT deposit money before "W" thread withdraws money.

## States of a Thread



25

## Two Important Things (1)

- You can safely change the state/value of a shared variable after calling `signalAll()`.
  - AS FAR AS the state changes in atomic code
- Common programming convention/practice:
  - A state change first, followed by `signalAll()`.

## Two Important Things (2)

- A JVM can perform context switches even when a thread runs atomic code.
  - A lock guarantees that only one thread exclusively runs atomic code at a time.
- Some resources (books, online materials, etc.) explicitly/implicitly say that context switches never occur when a thread runs atomic code.
  - It is WRONG!

## signal() and signalAll()

- `signalAll()`
  - Wakes up all waiting threads on a condition object.
    - All of them go to the “runnable” state.
    - One of them will re-acquire a lock. The others will go to the “waiting” state.
- `signal()`
  - Wakes up one of waiting threads on a condition object.
    - The selected thread goes to the “runnable” state. The others stay at the “waiting” state.
    - JVM’s thread scheduler selects one of them. Assume random selection.
      - Not predictable which waiting thread to be selected.

## signal() and signalAll()?

- Either one works well.
- signalAll() is favored in many cases/projects.
  - I prefer signalAll() in my personal taste.

## ThreadSafeBankAccount2.java

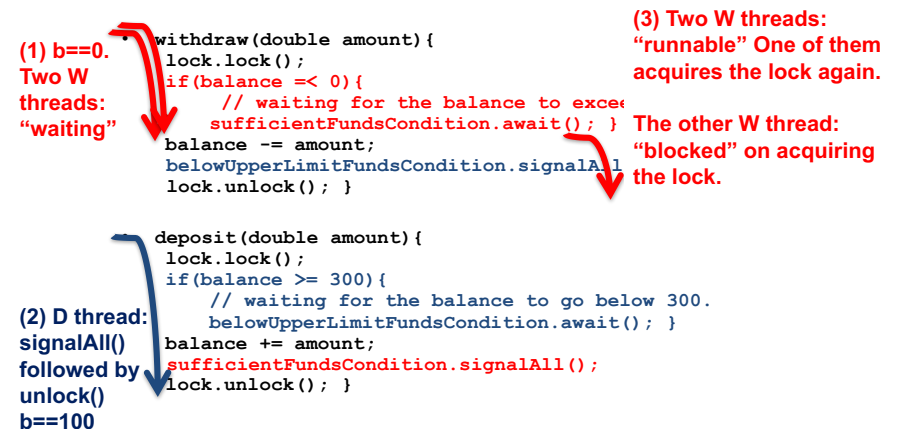
- Condition `sufficientFundsCondition` = lock.newCondition();  
Condition `belowUpperLimitFundsCondition` = lock.newCondition();
- `withdraw(double amount){`  
    `lock.lock();`  
    `while(balance <= 0){`  
        // waiting for the balance to exceed 0  
        `sufficientFundsCondition.await();` }  
    `balance -= amount;`  
    `belowUpperLimitFundsCondition.signalAll();`  
    `lock.unlock();` }
- `deposit(double amount){`  
    `lock.lock();`  
    `while(balance >= 300){`  
        // waiting for the balance to go below 300.  
        `belowUpperLimitFundsCondition.await();` }  
    `balance += amount;`  
    `sufficientFundsCondition.signalAll();`  
    `lock.unlock();` }

37

## “while” or “if” to Surround await()?

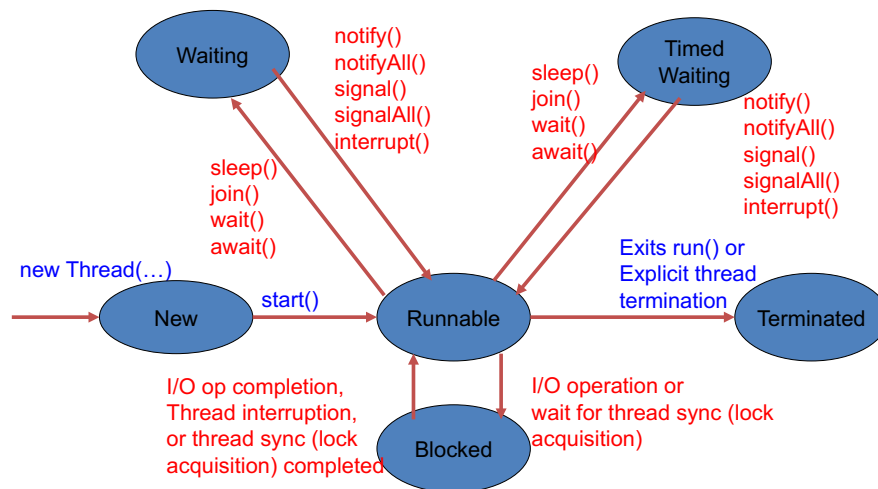
- `withdraw(double amount){`  
    `lock.lock();`  
    `while(balance <= 0){`  
        // waiting for the balance to exceed 0  
        `sufficientFundsCondition.await();` }  
    `balance -= amount;`  
    `belowUpperLimitFundsCondition.signalAll();`  
    `lock.unlock();` }
  - `withdraw(double amount){`  
    `lock.lock();`  
    `if(balance <= 0){`  
        // waiting for the balance to exceed 0  
        `sufficientFundsCondition.await();` }  
    `balance -= amount;`  
    `belowUpperLimitFundsCondition.signalAll();`  
    `lock.unlock();` }
- “while” should be used rather than “if” when multiple threads call withdraw() concurrently. Why?

## A Potential Problem

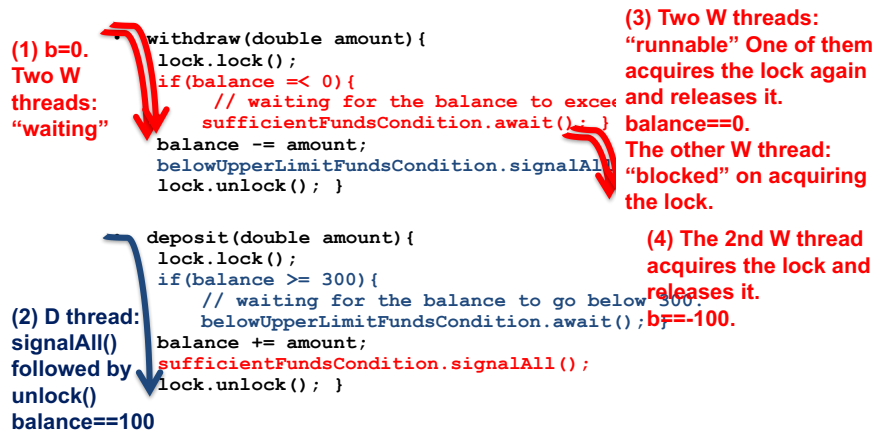
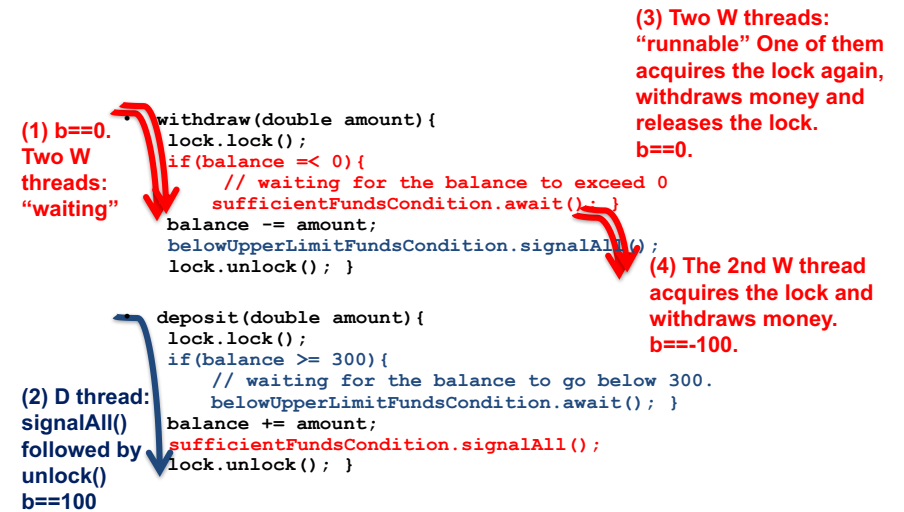




# States of a Thread



40



## "if" or "while" in Atomic Code?


- You can use "if", rather than "while," for a conditional check
  - if you use signal(), not signalAll().
- However, in practice, the **while-signalAll** pair is more common than the **if-signal** pair.

- The 2nd "W" thread should have made sure if  $balance > 0$ .
- If only one "W" thread runs, this problem does not occur.
- Just always use a while loop regardless of the number of threads you use.

## Producer-Consumer Design Pattern

## Producer-Consumer Design Pattern


- Producer
  - Generates data to be processed.
- Consumer
  - Take and process those generated data.
  - If no data is available, wait until a producer generates data.



45

## Bank Account Example

- DepositRunnable (Producer)
  - A thread (or a group of threads) that deposits money to a bank account.
  - If the current balance is over the upper limit, the thread(s) wait(s) until the balance goes below the upper limit.
- WithdrawRunnable (Consumer)
  - A thread (or a group of threads) that withdraws money from the account.
  - If the current balance is below the lower limit, the thread(s) wait(s) until the balance exceeds the lower limit.




## An Exercise: File Indexing Service

- Imagine an indexing service for a file system
  - e.g., Windows indexing service and Mac/iOS's Spotlight
- Key functionalities
  - Scan/crawl files in the local file system
  - Index those files for later file searches.
    - Extract and keep each file's metadata
      - Metadata: file's attributes (e.g., location, name, file type, author) and file's content

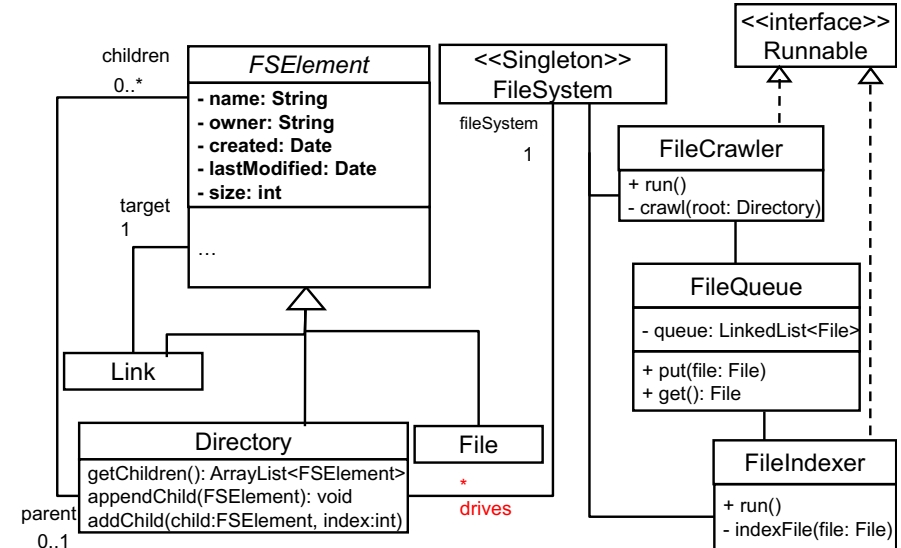
# Threads in File Indexing

- Single threaded
  - Use a single thread for both crawling and indexing
- Multi-threaded
  - Use different threads for crawling and indexing
    - One crawling thread and one indexing thread
    - Multiple crawling threads and multiple indexing threads
  - More efficient than the single-threaded version in multi-core environments
  - Crawlers and indexers interact with each other based on the Producer-Consumer design pattern.



- Assume multiple crawler threads and multiple indexer threads.
  - One crawler thread per a drive.
- A crawler thread
  - Traverses a tree structure in a given drive and puts files to the queue.
  - Waits, if the queue's size reaches a certain number (upper limit), until the size becomes below the limit.
  - Dies when it completes to traverse a given tree structure or when the main thread tells it to die.
- An indexer thread
  - Gets a file from the queue and indexes it.
    - Waits, if no files are available in the queue, until a crawler puts a new file.
  - Repeats this forever until the main thread tells it to die.

# HW 14: Implement this.



```

• class FileCrawler implements Runnable{
 private Directory root; //root dir of a given drive (tree structure)
 private FileQueue queue;
 ...
 public void run(){
 crawl(root);
 ...
 }
 private void crawl(Directory root){
 // Crawl a given drive (tree structure)
 // Put files to a queue. Ignore directories and links.
 queue.put(file);
 }
}

• class FileIndexer implements Runnable{
 private FileQueue queue;
 ...
 public void run(){
 while(true){
 indexFile(queue.get());
 }
 }
 public indexFile(File file){
 // Index a given file.
 }
}

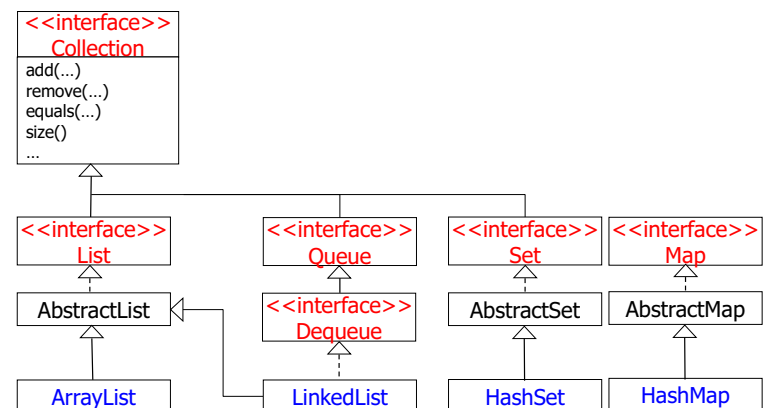



```

- No need to include a interactive command-line I/F
- No need to filter files in crawl()
  - crawl() can queue all files to the file queue
- No need to implement actual indexing logic.
  - indexFile() can just print out each file's metadata (e.g. file name) on a shell.
- Make multiple drives and assign a crawler thread to each drive.
- Run multiple crawler threads and multiple indexer threads from main().
- Have the main thread terminate crawler and indexer threads.
  - Use 2-step thread termination.

- Use LinkedList<File> in FileQueue.

## Just in Case: Major Collection Types in Java


- List
  - Orders elements with their integer **index numbers**.
  - Offers index-based **random** access.
  - Can contain **duplicate** elements.
- Queue
  - Orders elements with **links**.
  - Offers **FIFO** (First-In-First-Out) access.
  - Can contain **duplicate** elements.
- Dequeue
  - Stands for “Double Ended QUEUE” (pronounced “deck”).
  - Orders elements with **links**.
  - Offers both **FIFO and LIFO** (Last-In-First-Out) access.
  - Can contain **duplicate** elements.
- Set
  - Contains **non-duplicate** elements **without an order**.
- Map
  - Contains key-value pairs (w/ non-duplicate keys) **without an order**.



# ArrayList v.s. LinkedList


- ArrayList

- Array-based implementation of the List interface



- LinkedList


- Doubly-linked implementation of the List and Deque interfaces



57


- ArrayList

- Array-based impl of the List interface
- Fast index-based access
- Slow insertion and removal of non-tail elements
  - Fast insertion and removal of the tail element



- LinkedList


- Doubly-linked impl of the List and Deque interfaces



58


- ArrayList

- Array-based impl of the List interface
- Fast index-based access
- Slow insertion and removal of non-tail elements
  - Fast insertion and removal of the tail element



- LinkedList


- Doubly-linked impl of the List and Deque interfaces
- Fast insertion and removal of elements



59


- ArrayList

- Array-based impl of the List interface
- Fast index-based access
- Slow insertion and removal of non-tail elements
  - Fast insertion and removal of the tail element



- LinkedList


- Doubly-linked impl of the List and Deque interfaces
- Fast insertion and removal of elements



60


- ArrayList

- Array-based impl of the List interface
- Fast index-based access
- Slow insertion and removal of non-tail elements
  - Fast insertion and removal of the tail element



- LinkedList

- Doubly-linked impl of the List and Deque interfaces
- Fast insertion and removal of elements
- Slow index-based access for “middle” elements.



61

- Use ArrayList

- If you often need to access “middle” elements.

- Use LinkedList

- If you often need to insert/remove elements.

- Both yield the same performance for element traversal (i.e. sequential element access).

62

- Due: May 1 (Tue) midnight