

Immutable Objects

Immutability and Threading

- Objects that **never change their states**.
 - Only getter methods; no setter methods.
- **No need to worry about race conditions**.
 - It is always thread-safe.
 - No locking; no performance loss.
- An example: `java.lang.String`
 - `char[] str = {'u', 'm', 'b'};`
`String string = new String(str);`
 - `String string = "umb";`
 - A series of constructors to initialize string data.
 - All non-constructor methods never change the initialized string data.

2

Methods in String

```
• String str = "umb";  
    // Create a String instance that contains "umb"  
    // and assigns it to str.  
• System.out.println( str );           // umb  
• System.out.println( str.length() );   // 3  
• System.out.println( str );           // umb  
  
• System.out.println( str.replace("b","l")); // uml  
    // replace() creates a new String instance that  
    // contains "uml" and returns it.  
• System.out.println( str );           // umb  
    // str still contains "umb"  
  
• System.out.println( str.toUpperCase() ); // UMB  
    // toUpperCase() creates a new String instance  
    // that contains "UMB" and returns it.  
  
• System.out.println( str );           // umb  
    // str still contains "umb"
```

3

```
• String str = "umb";  
    // Create a String instance that contains "umb"  
    // and assigns it to str.  
• System.out.println( str );           // umb  
  
• str = "uml";  
    // Create another String instance that contains "uml"  
    // and assigns it to str.  
• System.out.println( str );           // uml
```

4

String

- Maintains the initialized string data (e.g. “umb”) in its *private* and *final* variable:
 - `private final char value[];`
 - Once a value is assigned to a final variable, the value cannot be changed afterward.
 - No methods of `String` can change the value.
- Is defined as a *final* class.
 - A final class cannot be extended (sub-classed).
 - Prevents its sub-classes from updating the initialized string data.

5

Note that...

- Whenever appropriate, use *both final and private* as often as possible in multi-threaded programs.
 - A small change can turn thread-safe code to be thread-unsafe.
- Clearly state immutability in program comments, API documents, design documents, etc.
 - Use {frozen} or {immutable} in UML class diagrams

7

Fragile Immutable Class

- ```
class Str{
 protected String data;

 StrData(String data){ this.data = data; }
 public String getStr(){ return data; } }
```
- ```
class Str2 extends Str{
    ChildStrData(String data){ super(data); }
    public void updateStr(String newStr){
        data = newStr; } }
```
- `Str` is immutable itself, but it is *fragile*.
 - *Immutable*: It has no public setter methods.
 - *Fragile*: It allows its subclasses to be mutable (i.e. to have public setter methods).
 - Should have been a *final class* that has a *private and final data field*.

6

Other Immutable Classes

- Wrapper classes for primitive types
 - `java.lang.Boolean` for boolean
 - `java.lang.Byte` for byte
 - `java.lang.Character` for char
 - `java.lang.Double` for double
 - `java.lang.Float` for float
 - `java.lang.Integer` for int
 - `java.lang.Long` for long
 - `java.lang.Short` for short
 - ...etc.

8

Immutable Classes are Good for...

- For API designers
 - Immutable classes never require locking (thread synchronization) for read operations.
 - They are free from concurrency issues.
 - Makes it easier to do debugging.
- for API users
 - Immutable classes are free from performance loss due to locking.

Well, Not All Classes can be Immutable...

- Immutable classes are useful.
- However, in practice, most classes are (or must be) mutable.
- Think of **separating a class to mutable and immutable parts**
 - if read operations are called very often.

9

An Example: String and StringBuilder

- Both represent string data.
- **String**
 - **Immutable**: Its state never change.
 - Thread-safe
 - Faster (or equally fast) to perform **read operations**.
- **StringBuilder**
 - **Mutable**: Its state can change through its public methods.
 - Not thread-safe
 - A LOT faster to perform **write operations**.

- ```
String str = "UMass";
str = str + " Boston"; // "UMass Boston"
// String concatenation. Automatically transformed to:
// str = new StringBuilder().append(str1).append(str2).toString();
// Creates two instances and calls two methods.
```
- ```
StringBuilder sbuilder(str);
sbuilder.append(" Boston");
str = sbuilder.toString();
```
- Write operations (incl. `append()`, `replace()`, etc.) can run **faster** with `StringBuilder`.

11

- ```
ArrayList<String> emailAddrs = ...;
String commaSeparatedEmailAddrs;
for(emailAddr: emailAddrs){
 commaSeparatedEmailAddrs += emailAddr + ", "; }
```
- ```
StringBuilder commaSeparatedEmailAddrs;
for(emailAddr: emailAddrs){
    commaSeparatedEmailAddrs.append(emailAddr).append(", "); }
```
- The latter code can run **20-100% faster** depending on the number of collection elements.

StringBuffer


- Provides the same set of public methods as `StringBuilder` does.
- `StringBuffer` (since Java 1.0)
 - All public methods are thread-safe with locking.
 - Client code of `StringBuffer` may still require locking.
- `StringBuilder` (since Java 5)
 - All public methods are NOT thread-safe.
 - Client code of `StringBuilder` require locking.

- Use `String` (immutable class) for read operations
- Use `StringBuilder` (mutable class) for write operations
- `String`-to-`StringBuilder` conversion is implemented in a constructor of `StringBuilder`.
- `StringBuilder`-to-`String` conversion is implemented in a constructor of `String`.

User-defined Immutable Class

Lock-ordering Deadlocks

Lock-ordering Deadlocks



```

class LockOrderingDeadlock{
    private ReentrantLock left = new ReentrantLock();
    private ReentrantLock right = new ReentrantLock();

    public void leftRight(){
        left.lock();
        right.lock();
        // atomic code
        right.unlock();
        left.unlock();
    }

    public void rightLeft(){
        right.lock();
        left.lock();
        // atomic code
        left.unlock();
        right.unlock();
    }
}

```

```

class LockOrderingDeadlock{
    private ReentrantLock left = new ReentrantLock();
    private ReentrantLock right = new ReentrantLock();

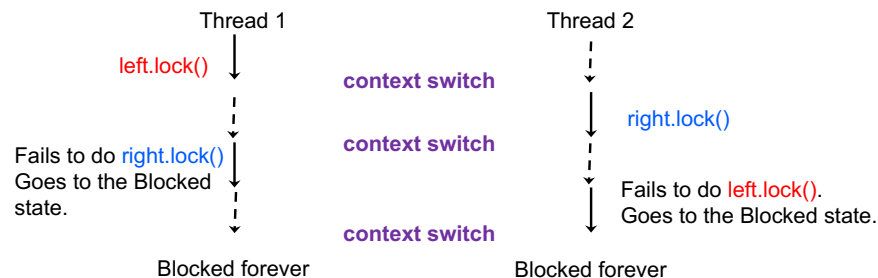
    public void leftRight(){
        left.lock();
        right.lock();
        // atomic code
        right.unlock();
        left.unlock();
    }

    public void rightLeft(){
        right.lock();
        left.lock();
        // atomic code
        left.unlock();
        right.unlock();
    }
}

```


A context switch can occur here.

A context switch can occur here.




- Problem:
 - Threads try to acquire *the same set of locks* in *different orders*.
 - Inconsistent lock ordering
 - Thread 1: left → right
 - Thread 2: right → left
- To-do:
 - Have all threads acquire the locks in a *globally-fixed* order.
- Be careful when you use multiple locks in order!

Dynamic Lock-ordering Deadlocks



```
class BankAccount{
    private ReentrantLock lock;
    ...
    public void deposit(...)
    public void withdraw(...)
    public void transfer(...)
}
```


- `void deposit(double amount){`
`lock.lock();`
`balance += amount;`
`lock.unlock(); }`
- `void withdraw(double amount){`
`lock.lock();`
`balance -= amount;`
`lock.unlock(); }`
- `void transfer(Account destination, double amount){`
`lock.lock();`
`if(balance < amount)`
 `// generate an error msg or throw an exception`
`else{`
 `withdraw(amount); // Nested locking. No problem.`
 `destination.deposit(amount); // Acquire another lock.`
`}`
`lock.unlock(); }`



```
class BankAccount{
    private ReentrantLock lock;
    ...
    public void deposit(...)
    public void withdraw(...)
    public void transfer(...)
}
```

- `void transfer(Account destination, double amount){`
`lock.lock();`
`if(balance < amount)`
 `// generate an error msg or throw an exception`
`else{`
 `withdraw(amount); // Nested locking. No problem.`
 `destination.deposit(amount); // Acquire another lock.`
`}`
`lock.unlock(); }`

- It looks as if all threads acquire the two locks (this.lock and destination.lock) in the same order.
- However, this code can have a lock-ordering deadlock.



Thread #1

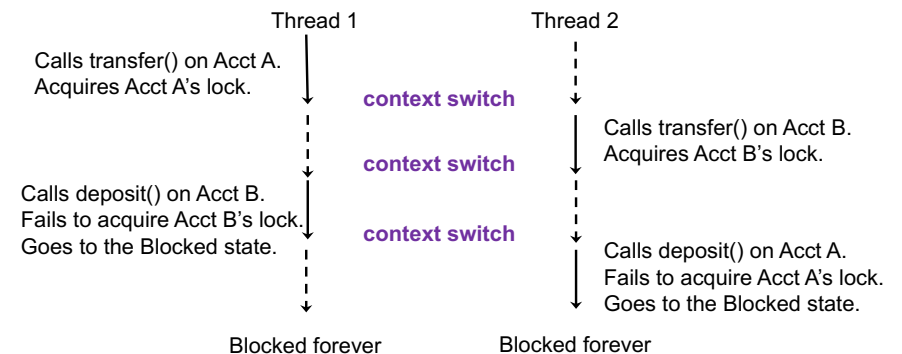
Threads can deadlock in transfer().

Thread #2

- Imagine a scenario where
 - a thread (#1) transfers money from Account A to B
 - another thread (#2) transfers money from B to A.

- `public void transfer(Account destination, double amount){`
`lock.lock();`
`if(balance < amount)`
 `// generate an error msg or...`
`else{`
 `withdraw(amount);`
 `destination.deposit(amount);`
`}`
`lock.unlock(); }`

A context switch can occur here.



Solutions

- Problem
 - Threads try to acquire *the same set of locks* in *different orders*.
 - *Inconsistent* lock ordering.
 - Thread 1: Acct A's lock → Acct B's lock
 - Thread 2: Acct B's lock → Acct A's lock
 - This can occur with bad timing although code looks OK.
 - A -> B and C -> D at the same time (No lock-ordering deadlock)
 - A -> B and B -> A at the same time (Possible lock-ordering deadlock)
- **Be careful when you use multiple locks in order!!!**

Solution 1: Static Lock

```
• private static ReentrantLock lock = new ReentrantLock();

• public void deposit(double amount){
    lock.lock();
    this.balance += amount;
    lock.unlock(); }

• public void transfer(Account destination, double amount){
    lock.lock();
    if( this.balance < amount )
        // generate an error msg or throw an exception
    else{
        this.withdraw(amount);      // Nested locking
        destination.deposit(amount); // Nested locking!
    }
    lock.unlock(); }
```

- Static lock
- Timed locking
- Ordered locking
- Nested tryLock()

- Pros
 - Simple solution
- Cons
 - Performance penalty
 - Transfers on different accounts are performed sequentially (not concurrently).
 - Deposit operations on different accounts are performed sequentially (not concurrently).
 - Withdrawal operations on different accounts are performed sequentially (not concurrently).

Solution 2: Timed Locking

```


• public void deposit(double amount){
    if( !lock.tryLock(3, TimeUnit.SECONDS) ){
        // generate an error msg or throw an exception
    }else{
        this.balance += amount;
        lock.unlock(); } }

• public void transfer(Account destination, double amount){
    lock.lock();
    if( this.balance < amount )
        // generate an error msg or throw an exception
    else{
        this.withdraw(amount); // Nested locking
        destination.deposit(amount);
    }
    lock.unlock(); } //Make sure to release this lock when
                        // an error/exception occurs.

```

- Pros
 - Simple solution
 - More efficient than Solution #1
 - By using a non-static lock
- Cons
 - Transfers and deposits might never be completed.
 - May look like unprofessional.

Solution 3: Ordered Locking



```

• public void transfer( Account source,
                        Account destination,
                        double amount){
    if( source.getAcctNum() < destination.getAcctNum() ){
        source.getLock().lock();
        destination.getLock().lock();
        if( source.getBalance() < amount )
            // generate an error msg or throw an exception
        else{
            source.withdraw(amount); //Nested locking
            destination.deposit(amount); //Nested locking
        }
        destination.getLock().unlock();
        source.getLock().unlock();
    }
    else if( source.getAcctNum() > destination.getAcctNum() ){
        destination.getLock().lock();
        source.getLock().lock();
        ...
        source.getLock().unlock();
        destination.getLock().unlock();
    }
}

```


Solution 3a: Ordered Locking with Instance IDs

- Pros
 - Locks are always acquired in the same order.
 - More efficient than Solution #1
 - By using a non-static lock
 - More professional than Solution #2
 - Transfers and deposits complete for sure.
- Cons
 - Using an application-specific/dependent data.
 - Account numbers should not be changed after accounts are set up.
 - If you allow dynamic changes of account numbers, you need to use an extra lock.


```
• public void transfer( Account source,
                      Account destination,
                      double amount){
    int sourceID = System.identifyHashCode(source);
    int destID = System.identifyHashCode(destination);

    if( sourceID < destID ){
        source.getLock().lock();
        destination.getLock().lock();
        if( source.getBalance() < amount )
            // generate an error msg or throw an exception
        else{
            source.withdraw(amount); //Nested locking
            destination.deposit(amount); //Nested locking
        }
        destination.getLock().unlock();
        source.getLock().unlock();
    }
    if( sourceID > destID ){
        destination.getLock().lock();
        source.getLock().lock();
        ...
        source.getLock().unlock();
        destination.getLock().unlock();
    }
}
```

- Instance IDs
 - Unique IDs (hash code) that the local JVM assigns to individual class instances.
 - Unique and intact on the same JVM
 - 2 instances of the same class have different IDs.
 - No instances share the same ID.
 - IDs never change after they are assigned to instances.
 - Use `System.identifyHashCode()`

- Pros
 - Locks are always acquired in the same order.
 - More efficient than Solution #1
 - By using a non-static lock
 - More professional than Solution #2
 - Transfers and deposits complete for sure.
 - No application-specific data (e.g., account numbers) are necessary to order locking.
- Cons
 - N/A

Solution 4: Nested Timed Locking



- Use nested tryLock() calls to implement an ALL-OR-NOTHING policy.
 - Acquire both of A's and B's locks, OR
 - Acquire none of them.
- Avoid a situation where a thread acquires one of the two locks and fails to acquire the other.

```

• public void transfer(Account source,
    Account destination,
    double amount){
    Random random = new Random();

    while(true){
        if( source.getLock().tryLock() ){
            try{
                if( destination.getLock().tryLock() ){
                    try{
                        if( source.getBalance() < amount )
                            // generate an error msg/exception
                        else{
                            source.withdraw(amount);
                            destination.deposit(amount);
                        }
                        return;
                    }finally{
                        destination.getLock().unlock();
                    }
                }
            }finally{
                source.getLock().unlock();
            }
        }
        Thread.sleep(random.nextInt(1000));
    }
}

```

- If the first tryLock() fails, then sleep.
- If the first tryLock() succeeds but the second one fails, unlock the first lock and sleep.

- Pros
 - More efficient than Solution #1
 - By using a non-static lock
 - More professional than Solution #2
 - Transfers and deposits complete for sure.
 - No application-specific data (e.g., account numbers) are necessary to order locking.
- Cons
 - Not that simple