

## Collections and Concurrency

### 3 Types of Collections in Java

- Thread-unsafe collections
- Thread-safe collections
  - Synchronized collections
  - Concurrent collections

### Thread-unsafe Collections

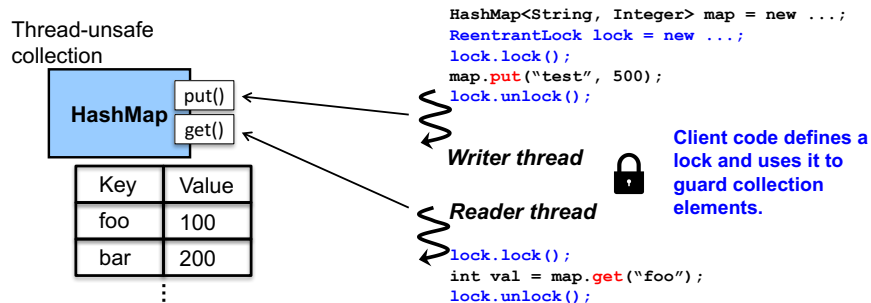
- Many collection classes are **NOT thread safe**.
  - e.g., `ArrayList`, `LinkedList`, `HashMap`, etc.
  - Their public methods **never** perform thread synchronization (i.e., locking).
- Look into Java API documentation to see if a collection is thread-safe or not.

### Java API Doc on `ArrayList`

- “Note that this implementation is **not synchronized**. If multiple threads access an `ArrayList` instance concurrently, and at least one of the threads modifies the list structurally, it *must* be synchronized externally. (A structural modification is any operation that adds or deletes one or more elements, or explicitly resizes the backing array; merely setting the value of an element is not a structural modification.)...”

## When You Use a Thread-unsafe Collection...

- You must do *client-side locking*; your client code must perform *thread synchronization* (i.e., locking)
  - to guard collection elements against concurrent accesses.
    - c.f. previous HWs that use thread-unsafe collections such as `ArrayList`, `LinkedList` and `HashMap`.



## Example Compound Operations

- ```

List<String> list =
    Collections.synchronizedList( new ArrayList<String>());

Iterator it = list.iterator();
while( it.hasNext() ){           // Iteration
    doSomething( it.next() ); }
    
```
- ```

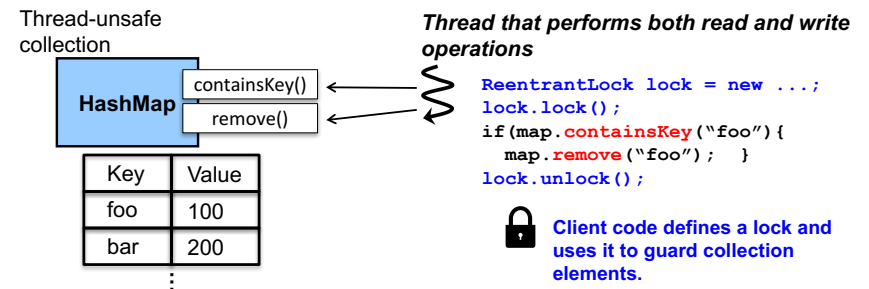
for(int i=0; i<list.size(); i+=2){ // Navigation
    doSomething( list.get(i) ); }
    
```
- ```

if( list.size() > 10 )           // Check-then-act
    doSomething( list );
    
```

- You must do *client-side locking* for *compound operations* as well as simple public method calls.
- Example compound operations
  - Iteration (element-traversal)*
    - Repeatedly get elements one by one until a collection becomes empty
  - Navigation*
    - Find/search the next element after a given element
  - Conditional operations (check-then-act)*
    - e.g., Check if a `HashMap` has a key-value pair for the key `k`, and if not, add the pair `(k,v)`

## Client-side Locking for Compound Operations

- An example “check-then-act” compound operation



# Synchronized Collections

- “Ready-made” **thread-safe** collections
  - **Synchronized classes**: `Vector` and `Hashtable`
    - All public methods perform **thread synchronization** (locking).
      - Only one thread can access the elements of a collection at a time.
        - » e.g., When a thread is in the middle of executing `add()` on a `Vector`, no other threads can call `get()`, `size()`, etc. on that `Vector`.
  - **Synchronized wrapper classes** for thread-unsafe collections
    - Created by `java.util.Collections.synchronizedXyz()`
      - Factory methods
        - `synchronizedList()`
        - `synchronizedMap()`
        - `synchronizedSet()`

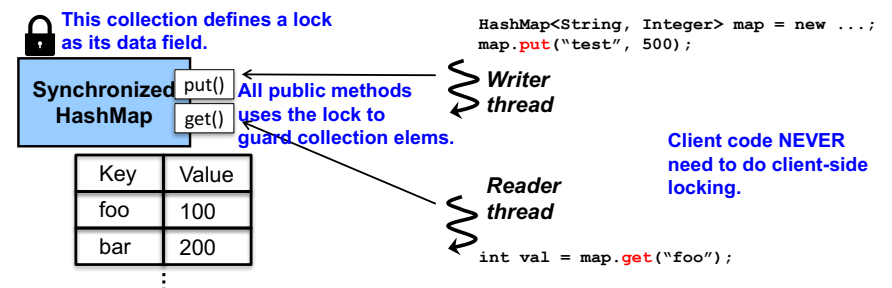
## Synchronized Wrapper Classes

- `List<String> list = Collections.synchronizedList( new ArrayList<String>() );`
- `list`: an instance of a synchronized wrapper class for `ArrayList`.
  - `list.getClass()` → `java.util.Collections$SynchronizedRandomAccessList`
- The wrapper class offers “synchronized” (or **thread-safe**) versions of `ArrayList`’s public methods.
  - `add()`, `get()`, `remove()`, etc.

## Vector and Hashtable in Single Threaded Programs

- It makes no sense to use these collections in single-threaded programs in a **performance point of view**.
  - They perform thread synchronization (locking) even when only one thread runs in a program.
    - **Unnecessary performance loss**
- Use **`ArrayList`** and **`HashMap`** instead!

- All public methods are **thread-safe** in all synchronized collection classes.
- No client-side locking is necessary when client code makes a simple public method call.



## Potential Problems in Compound Actions

- Need *client-side locking in compound operations* on a synchronized collection
  - Iteration (element-traversal)
    - Repeatedly get elements one by one until a collection becomes empty
  - Navigation
    - Find/search the next element after a given element
  - Conditional operations (check-then-act)
    - e.g., Check if a `HashMap` has a key-value pair for the key `k`, and if not, add the pair `(k,v)`

```
• List<String> list =  
  Collections.synchronizedList( new ArrayList<String>());  
  
Iterator it = list.iterator();  
while( it.hasNext() )           // Iteration  
  doSomething( it.next() );  
  
• if( list.size() > 10 )         // Check-then-act  
  doSomething( list );  
                                Race conditions can occur here.
```

- Race conditions
- `ConcurrentModificationException`
  - Raised if a writer thread tries to add/remove elements before a reader thread completes a traversal on the entire set of elements.

## Client-side Locking for Compound Operations

- ```
• List<String> list =  
  Collections.synchronizedList( new ArrayList<String>());  
  
• synchronized(list){  
  Iterator it = list.iterator();  
  while( it.hasNext() )           // nested locking  
    doSomething( it.next() );     // nested locking  
}  
  
• synchronized(list){  
  if( list.size() > 10 )           // nested locking  
    doSomething( ... );  
}  
  
• synchronized(list): acquires the lock that the list owns/uses for thread synchronization in its public methods.
```

- ```
• ReentrantLock lock = new ReentrantLock();  
lock.lock();                       // acquires lock  
Iterator it = list.iterator();  
while( it.hasNext() )              // acquires the lock that the  
  doSomething( it.next() );         // list owns.  
lock.unlock();                     // nested locking
```
- `lock` is different from the lock that the `list` owns/uses for thread synch in its public methods.
    - Must make sure to use `lock` consistently in all client code of `list`.
  - A thread acquires 2 locks.

# Performance Implication on Client-side Locking

- ```
List<String> list = Collections.synchronizedList( new ArrayList<String>());
```
- ```
synchronized(list){
    Iterator it = list.iterator();
    while( it.hasNext() )           // nested locking
        doSomething( it.next() ); // nested locking
}
```
- ```
ReentrantLock lock = new ReentrantLock();
lock.lock();           // acquires lock
Iterator it = list.iterator();
while( it.hasNext() ) // acquires the lock that the
                        // list owns.
    doSomething( it.next() ); // nested locking
lock.unlock();
```
- Performance penalty due to 2 lock acquisitions
  - Whether it is nested locking or not.

## Summary: Thread un-safe Collections v.s. Synchronized Collections

- Thread un-safe collections
  - Client code always must perform locking (thread synchronization).
    - For **both simple public method calls and compound operations**
  - Client-side locking always requires 1 lock acquisition.
- Synchronized collections
  - Client code does NOT have to perform locking **for simple public method calls**.
    - This client-side locking require 1 lock acquisition.
  - However, it needs client-side locking **for compound operations**.
    - This client-side locking requires 2 lock acquisitions.

```
List<String> list = new ArrayList<String>();
...
ReentrantLock lock = new ReentrantLock();
lock.lock();
Iterator it = list.iterator();
while( it.hasNext() )
    doSomething( it.next() );
lock.lock();
```

- Performs client-side locking for **ArrayList**, which is **NOT** thread-safe.
  - The client code is thread-safe although `hasNext()` and `next()` are not.
- **More efficient** than the previous client code.
  - No 2 lock acquisitions (only 1 lock acquisition)
  - Use a read-write lock if you have many “reader” threads.

## Concurrent Collections

- “Ready-made” **thread-safe** collections
  - Introduced in Java 5 (2004) and enhanced in subsequent versions
    - Queue
      - `ConcurrentLinkedQueue` (since Java 5)
      - `ConcurrentLinkedDeque` (since Java 7)
      - `ArrayBlockingQueue` (since Java 5)
      - `LinkedBlockingQueue` (since Java 5)
      - `DelayQueue` (since Java 5)
      - `PriorirtyBlockingQueue` (since Java 5)
      - `LinkedBlockingDeque` (since Java 6)
      - `LinkedTransferQueue` (since Java 7)
    - Map
      - `ConcurrentHashMap` (since Java 5)
      - `ConcurrentSkipListMap` (since Java 6)
    - Set
      - `ConcurrentSkipListSet` (since Java 6)
  - `java.util.concurrent.CopyOnWriteXyz` classes
    - `CopyOnWriteArrayList`
    - `CopyOnWriteArraySet`

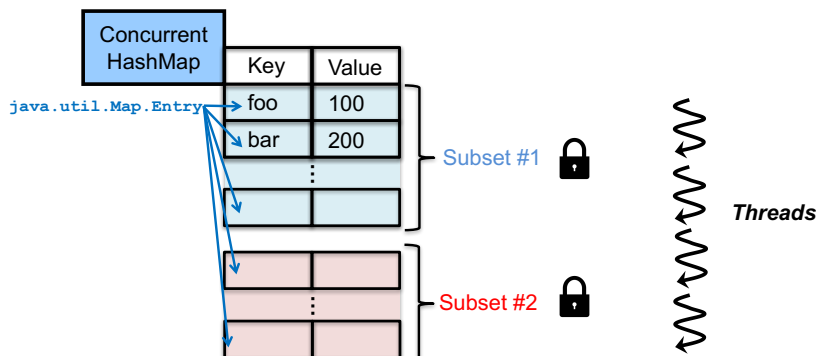
# What is ConcurrentHashMap?

- Indented to **replace synchronized collections**.
  - e.g., `ConcurrentHashMap` is indented to replace `SynchronizedMap`.
- Make public methods **thread-safe**.
  - e.g., `get()` and `put()` in `ConcurrentHashMap`
  - Client code does NOT have to perform locking **for simple public method calls**.
- Implement public methods in an **efficient** manner.
  - Perform lock stripping
- Provide **thread-safe methods to perform compound operations**
  - Client code does not have to perform locking **for many compound operations**.

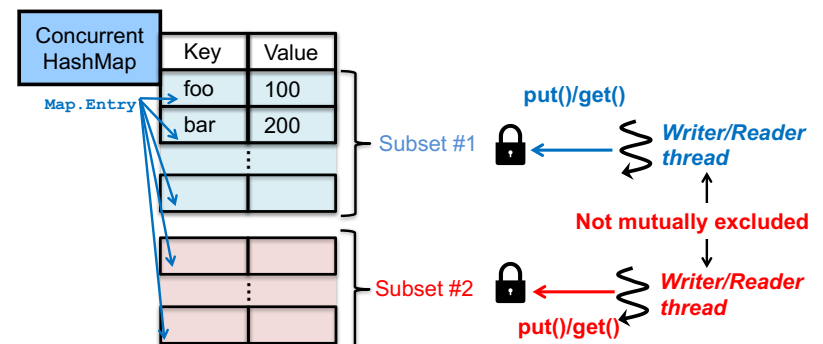
- Provides thread-safe public methods.
  - e.g., `get()` and `put()`
  - Client code does NOT have to perform locking **for simple public method calls**.
- A replacement for synchronized hash-based `Map` implementations (e.g., `Hashtable` and synchronized `HashMap`).
- Implement public methods in an **efficient** manner.
  - Performs ***fine-grained locking***, called **lock stripping**
    - Compared to *coarse-grained* locking (i.e., lock-per-collection) in synchronized hash-based `Map` implementations.

## Lock Stripping

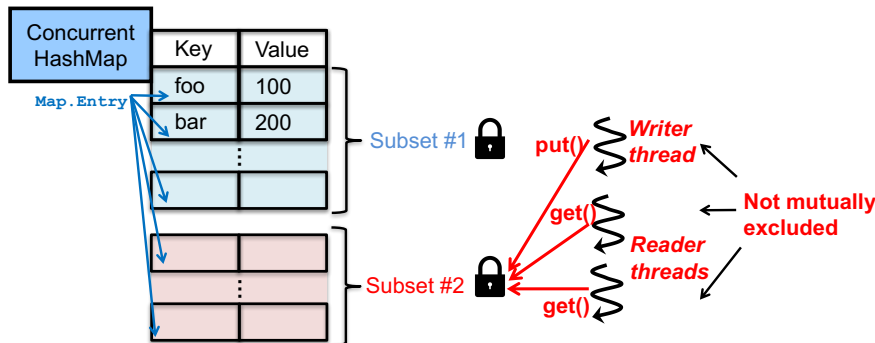
- `ConcurrentHashMap` uses **multiple** locks to guard a table (i.e., key-value pairs).
  - 16 locks by default
    - Configurable with the “concurrencyLevel” parameter in a constructor.
  - Each lock is associated with a subset of the table.



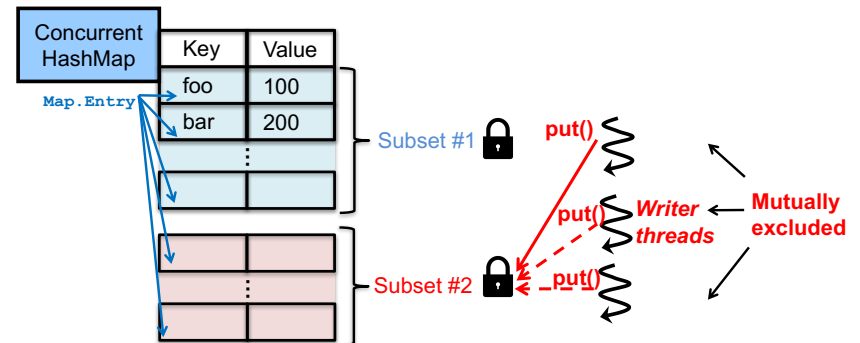
- Threads are **not synchronized (i.e., not mutually excluded)** with each other
  - as far as they access different subsets of the table with different locks.



- To access a subset of the table,
  - Reader threads
    - are **NOT synchronized (NOT mutually excluded)** with each other.
      - c.f. read-write lock
    - are **NOT synchronized (NOT mutually excluded)** with writer threads.
      - c.f. inner class Node<K, V>

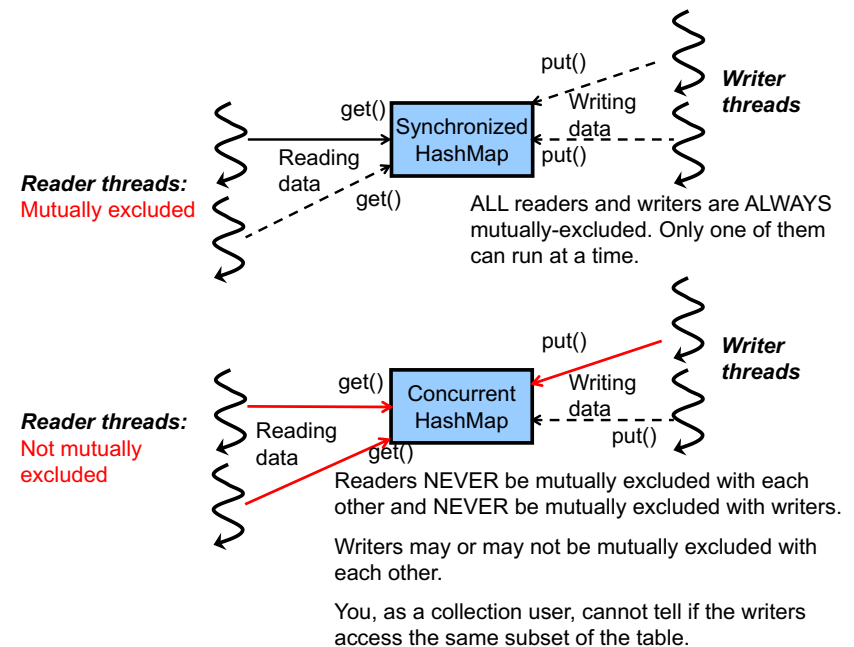



- To access a subset of the table,
  - Writer threads **ARE synchronized** (mutually excluded) with each other.



## Synchronized Hash-based Map Impls

- A **single lock** is used to guard the *entire* table in `Hashtable` and synchronized `HashMap`
  - No lock stripping.
- All writer/reader threads ARE ALWAYS synchronized (mutually excluded) with each other.
  - A potential performance bottleneck as the number of key-value pairs increases and the number of threads increases.




## Thread-safe Iteration with Concurrent Iterators

- Supports thread-safe iteration
  - Does not require client-side locking
  - c.f. Thread-unsafe collections and synchronized collections require client-side locking for iteration
    - Because it is a compound operation.

```
• Iterator it = aConcurrentHashMap.entrySet().iterator();
  while( it.hasNext() )
    doSomething( it.next() );
```

### – Pros

- No client-side locking is necessary in client code.
  - Readers are fully concurrent (not mutually excluded).
  - Writers can add/remove elements while readers read elements.
    - » It is guaranteed that writers and readers do not corrupt elements.
- The iterator “it” is *backed by* the map, so changes to the map are reflected in the set.




## Concurrent Iterators

- Concurrent iterators are obtained through `entrySet()`, `keySet()` and `values()`.
  - `entrySet()`: Returns key-value pairs as a Set.
    - `Set<Map.Entry<K,V>>`
  - `keySet()`: Returns keys as a Set.
    - `ConcurrentHashMap.KeySetView<K,V>`
  - `values()`: Returns values as a Collection.
    - `Collection<V>`
- `Iterator it = aConcurrentHashMap.entrySet().iterator();`  
`while( it.hasNext() )`  
`doSomething( it.next() );`

```
• Iterator it = aConcurrentHashMap.entrySet().iterator();
  while( it.hasNext() )
    doSomething( it.next() );
```

### – Pros

- No client-side locking is necessary in client code.
  - Readers are fully concurrent (not mutually excluded).
  - Writers can add/remove elements while readers read elements.
    - » It is guaranteed that writers and readers do not corrupt elements.
- The iterator “it” is *backed by* the map, so changes to the map are reflected in the set.





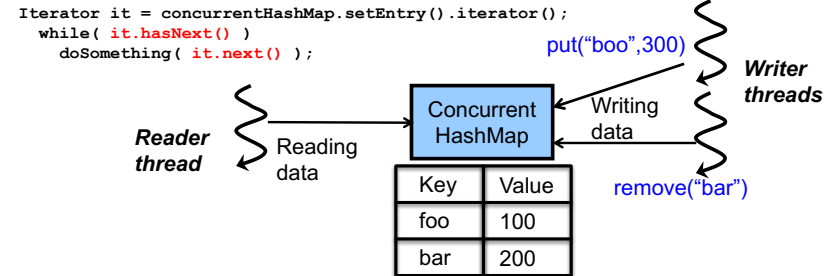
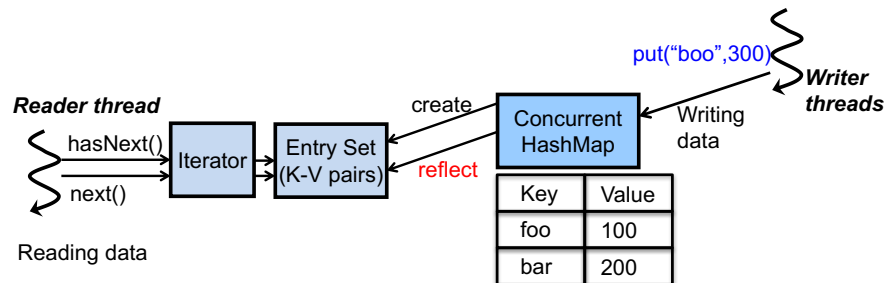
```

• Iterator it = aConcurrentHashMap.entrySet().iterator();
while( it.hasNext() )
    doSomething( it.next() );

```

### – Cons: weak (or best-effort) consistency

- There is no guarantee about how soon a change (to the hash map) to be reflected into the entry set.
- The iterator “it” may or may not traverse the most up-to-date key-value pairs in the map.



- The iterator “it” may traverse
  - {(foo,100), (bar,200)},
  - {(foo,100), (bar,200), (boo,300)},
  - {(foo,100), (boo,300)}, or
  - {(foo,100)}.
- Elements are guaranteed not to get corrupted.
  - e.g., {(foo,300)} is never put to the map.

## Consistency v.s. Performance

- **ConcurrentHashMap** trades *perfect consistency* for *performance improvement*.
  - If you can live with weak consistency, it's a great collection class for you.
    - Pros: performance improvement
    - Cons:
      - » Iterators may or may not traverse the most up-to-date key-value pairs in the map.
      - » `mappingCount()` and `isEmpty()` are not perfectly reliable.
        - The value returned is an estimate; the actual value may differ if there are concurrent insertions or removals.
  - If you cannot, craft your own thread-safe hash map with **HashMap** and **ReentrantLock**.
    - **ConcurrentHashMap** has no built-in scheme to lock the entire map.
- **int size()**
  - Returns the total number of key-value pairs with `int`.
    - `int`: 32-bit signed integer: - 2147483647 to 2,147,483,647
    - What if you want/need to have more than 2.15 billion pairs?
- **long mappingCount()** [since Java 8]
  - Returns the total number of key-value pairs with `long`.
    - `Long`: 64-bit signed integer: -9223372036854775808 to 9,223,372,036,854,775,807
  - Use this method, rather than `size()`, when you maintain a huge number of key-value pairs.

## A New Method in Iterator

- `java.util.Iterator<E>`
  - Used to traverse individual elements in a collection
    - `Iterator it = concurrentHashMap.setEntry().iterator();`  
`while( it.hasNext() )`  
`doSomething( it.next() );`
- `forEachRemaining()` [since Java 8]
  - Accepts a lambda expression (LE) and applies the LE to each remaining element
    - until all elements have been processed or the action throws an exception.
  - `Iterator it = concurrentHashMap.setEntry().iterator();`  
`it.forEachRemaining( (Map.Entry<String, AtomicInteger> e)`  
`->{System.out.println(e); } );`
  - `Iterator it = concurrentHashMap.keySet().iterator();`  
`it.forEachRemaining( (String k)`  
`->{System.out.println(k); } );`
- No client-side locking is required to call `forEachRemaining()`

## Thread-safe Compound Actions


- Supports *common compound actions* in a thread-safe way
  - put-if-absent: `putIfAbsent(key, value)`
    - Insert a pair of `key` and `value` as a new entry if `key` is not already associated with a value.
  - Conditional remove: `remove(key, value)`
    - Remove the entry for `key` if `key` is associated with `value`.
  - Conditional replace: `replace(key, value)`
    - Replace the entry for `key` if `key` is associated with some value.
  - Conditional replace: `replace(key, oldValue, newValue)`
    - Replace the entry for `key` with `newValue` only if `key` is associated with `oldValue`.
  - No client-side locking is necessary
  - c.f., `ConcurrentMap` interface

## A Specialized Iterator: Splitter

- `java.util.Spliterator<E>`
  - “Split” + “iterator”
  - Can *split* (or subset) the entire set of elements and traverse a subset of the elements.
- `Iterator it = concurrentHashMap.keySet().iterator();`  
`it.forEachRemaining( (String k)`  
`->{System.out.println(k); } );`
- `Spliterator st = concurrentHashMap.keySet().spliterator();`  
`Spliterator st2 = st.trySplit();`  
`st.forEachRemaining( (String k)`  
`->{System.out.println(k); } );`  
`st2.forEachRemaining( (String k)`  
`->{System.out.println(k); } );`
- `st` and `st2` cover two different subsets.

## Exercise: Concurrent Access Counter

- **AccessCounter**
  - c.f. HW 12 (w/ `HashMap` and `ReentrantLock`) and HW 15 (w/ `HashMap` and `ReentrantReadWriteLock`)
  - Maintains a map that pairs a relative file path and its access count.
  - `increment()`
    - accepts a file path and increments the file’s access count.
  - `getCount()`
    - accepts a file path and returns the file’s access count.



## Thread-safe Methods that Accept Lambda Expressions

- Imagine a CHM-based access counter

- `increment()`
  - ```
Integer oldCount = map.get(aFilePath);
int newCount = oldCount==null? 1: oldCount+1; // unboxing
map.put(aFileName, newCount);                // autoboxing
```

- This client code is not thread-safe.
  - because it performs a compound operation.
  - although `get()` and `put()` are thread-safe.
  - Need client-side locking

- Alternative: Use `putIfAbsent()` and `AtomicInteger`

- ```
ConcurrentHashMap<String, AtomicInteger> map = new ...;
map.putIfAbsent(aFilePath, new AtomicInteger(0));
map.get(aFileName).incrementAndGet();
```

## HW 19

- Revise your concurrent access counter (HW 12 solution) with `ConcurrentHashMap`
  - Eliminate client-locking (to guard the map) in `increment()` and `getCount()`
    - i.e., `AccessCounter` no longer has a `ReentrantLock` data field.
  - Use lambda expressions whenever possible.

- `compute()`
  - Accepts a key and a function (as a lambda expression).
  - Applies the function on a pair of the specified key and its value.
    - If the pair does not exist, the value contains null.
  - The function (re)maps the key and value.
  - ```
map.compute(aFilePath, (String k, Integer v)->
    {return v==null? 1: ++v;})
```
- `computeIfAbsent()`
  - ```
map.computeIfAbsent(aFilePath, (String k)->{return 1;})
```
- `computeIfPresent()`
  - ```
map.computeIfPresent(aFilePath, (String k, int v)->
    {return ++v;})
```

## Bulk Operations

- Repeatedly apply a given lambda expression on key-value pairs.
  - *forEach*: `forEach()`, `forEachEntry()`, `forEachKey()`, `forEachValue()`
  - *Search*: `search()`, `searchEntries()`, `searchKeys()`, `searchValues()`
  - *Reduce*:
    - `reduce()`, `reduceToDouble()`, `reduceToInt()`, `reduceToLong()`
    - `reduceEntries()`, `reduceEntriesToDouble()`, `reduceEntriesToInt()`, `reduceEntriesToLong()`
    - `reduceKeys()`, `reduceKeysToDouble()`, `reduceKeysToInt()`, `reduceKeysToLong()`
    - `reduceValues()`, `reduceValuesToDouble()`, `reduceValuesToInt()`, `reduceValuesToLong()`
- Receives a “parallelism threshold” as the first parameter.
  - Executed with threads if the # of key-value pairs exceeds this threshold.
    - `Long.MAX_VALUE`: Suppress concurrency/parallelism. Use a single thread.
    - 1: Maximize concurrency/parallelism.

## forEach Bulk Operations

- The number of threads to be used:
  - If (# of key-value pairs) < threshold
    - No threads to be used
  - If (# of KV pairs) / threshold >= (# of CPU cores)
    - Use (# of CPU cores)
  - If (# of KV pairs) / threshold < (# of CPU cores)
    - Use (# of KV pairs) / threshold
      - Rounded to an int number

- Performs a given action on each key-value pair.

```
map.forEach( 50,
    (k,v) -> {if (v>10000)
                System.out.println(k + "->" + v) } )
```

- First param: parallelism threshold (50)
- Second param: action (`BiConsumer`)
- Printing out the key-value pairs that have more than 10,000 access count.

## search Bulk Operations


- Searches a key-value pair that satisfies a given search criterion, which is specified as a lambda expression
  - Returns a non-null result if found. Returns null otherwise.
  - Skips further search when a result is returned.
- Path `path = map.search(50, (k,v) -> {v>10000? v: null} )`
  - Second param: search function (`BiFunction`)
    - Is there at least one file that has more than 10,000 access count?
  - If there exists such a file, `path` contains a path to that file.
  - Only the first search hit is returned. (Extra search hits are ignored)
- Integer `count = map.searchValues(50, (v) -> {v>10000? v: null} )`
  - Second param: search function (`Function`)
    - Is there at least one file that has more than 10,000 access count?
  - If there exists such a file, `count` contains the access count of that file.

## reduce Bulk Operations

- Reduces key-value pairs to a single value.

```
Int maxCount = map.reduceToInt(50,
    (k,v) -> v,
    0,
    (max, v) -> v>max? v: max );
```

- Second param: transformer (`BiFunction`)
  - Expected to work like a map operation
- Fourth param: reducer (`BinaryOperator`)
- Third param: the initial value for accumulated value (`max`).



# Notes on Bulk Operations

- `int maxCount = map.reduceToInt(50, (k,v)->v, 0, (max, v)-> v>max? v: max );`
  - Second param: transformer (`BiFunction`)
    - Expected to work like a map operation
  - Fourth param: reducer (`BinaryOperator`)
  - Third param: the initial value for accumulated value (`max`).
- Generalized form of reduce operations with the Streams API
  - `T result = aStream.reduce(initValue, (result, element)-> ... );`
  - `T result = initValue;`  
`for(T element: collection){`  
`result = accumulate(result, element); }`

## Recap: Concurrent Collections

- “Ready-made” **thread-safe** collections
  - Introduced in Java 5 (2004) and enhanced in subsequent versions
    - **Queue**
      - `ConcurrentLinkedQueue` (since Java 5)
      - `ConcurrentLinkedDeque` (since Java 7)
      - `ArrayBlockingQueue` (since Java 5)
      - `LinkedBlockingQueue` (since Java 5)
      - `DelayQueue` (since Java 5)
      - `PriorirtyBlockingQueue` (since Java 5)
      - `LinkedBlockingDeque` (since Java 6)
      - `LinkedTransferQueue` (since Java 7)
    - **Map**
      - `ConcurrentHashMap` (since Java 5)
      - `ConcurrentSkipListMap` (since Java 6)
    - **Set**
      - `ConcurrentSkipListSet` (since Java 6)
- `java.util.concurrent.CopyOnWriteXyz` classes
  - `CopyOnWriteArrayList`
  - `CopyOnWriteArraySet`

- Bulk operations
  - are **thread-safe** in that they never corrupt key-value pairs.
  - may not be performed on the most up-to-date key-value pairs.
    - Write threads can modify key-value pairs when read threads are reading key-value pairs.
      - It is guaranteed that read and write threads do not corrupt key-value pairs.
      - c.f. Concurrent iteration
    - Read threads can be fully concurrent as if a read-write lock is used in a CHM.

- Every concurrent collection...
  - Implements its public methods in **thread-safe** and **performance-aware** manners, just like `ConcurrentHashMap`.
  - Supports **thread-safe**, **weak-consistent** iteration.
    - May or may not provide thread-safe methods tailored for the other types of compound operations
      - c.f. `putIfAbsent(key, value)` and `compute()` in `ConcurrentHashMap`

## Some Major Concurrent Collections

- **ConcurrentLinkedQueue**
  - Concurrent implementation of `java.util.Queue`
    - FIFO (First-In-First-Out) queue
- **ConcurrentLinkedDeque**
  - Concurrent implementation of `java.util.Deque`
    - LIFO (Last-In-First-Out)
- **ConcurrentSkipListMap**
  - An implementation of `ConcurrentNavigableMap`.
  - Map entries are kept sorted according to the natural ordering of their keys or by a custom `Comparator`.
- **ConcurrentSkipListSet**
  - A concurrent implementation of `NavigableSet`.
  - Set elements are kept sorted according to the natural ordering or by a custom `Comparator`.

## Copy-On-Write (COW) Collections

- `CopyOnWriteArrayList`
- `CopyOnWriteArraySet`
  - Concurrent replacements of synchronized wrappers for `ArrayList` and `ArraySet`.
- Key features
  - Thread-safe public methods
    - No client-side locking is necessary to call them.
  - Read threads are never mutually excluded with each other.
  - Read threads and write threads are never mutually excluded.
  - Write threads are never mutually-excluded.


## Concurrent Hash Set???

- No concurrent collection class for **hash-based set** in Java API
  - No such class like `ConcurrentHashSet`.
- You can “emulate” it through `ConcurrentHashMap`:
  - `Set<String> set = ConcurrentHashMap<String, Integer>.newKeyset();`
    - Values (`Integers`) are ignored.
    - `set` is a **hash-based set** that contains a series of `String` data.
    - `newKeyset()` is a static factory method to generate a `ConcurrentHashMap.KeySetView<K, Boolean>`
      - `ConcurrentHashMap.KeySetView<String, Boolean>` in the above example.

## Snapshot-based Iteration in COW Collections

- Support thread-safe, **snapshot-based iteration**.
  - A read thread references and operates on a **collection snapshot**, which is a collection that is up-to-date when an iterator is created.


```
Iterator it = aCOWArrayList.iterator(); // A snapshot is created.
while( it.hasNext() )
    doSomething( it.next() );
```



- Pros

- No client-side locking is necessary for iteration.
  - Read threads ARE NOT mutually excluded with each other and with writer threads.
    - Each iterator has a thread-specific snapshot of List elements.
    - Different readers get different snapshots and access them concurrently.


```
Iterator it = aCOWArrayList.iterator(); // A snapshot is created.
while( it.hasNext() )
    doSomething( it.next() );
```



- Cons


- The snapshot can become **outdated**.
  - e.g., when a write thread adds/removes collection elements after a snapshot is created.
- No consistency preserved; Each iterator will **NOT** reflect additions, removals and other changes to the list after the iterator was created.
  - c.f. weak consistency in `concurrentHashMap` (and other non-COW concurrent collections)

- Trades consistency for performance





- Pros

- No client-side locking is necessary for iteration.
  - Read threads ARE NOT mutually excluded with each other and with writer threads.
    - Each iterator has a thread-specific snapshot of List elements.
    - Different readers get different snapshots and access them concurrently.



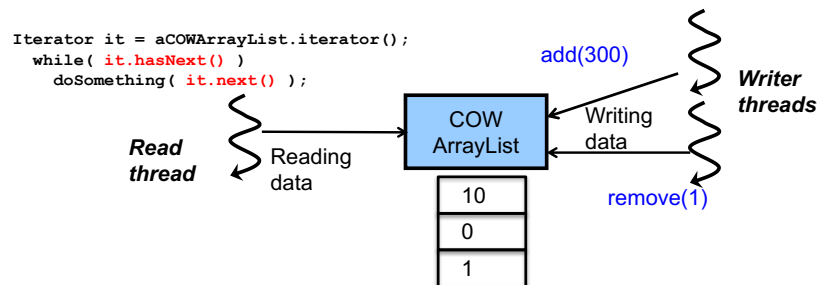

- Cons

- No consistency preserved. **Why?**
  - State updates (particularly, element insertion and removal) are often expensive for lists.
    - Due to index-based element ordering

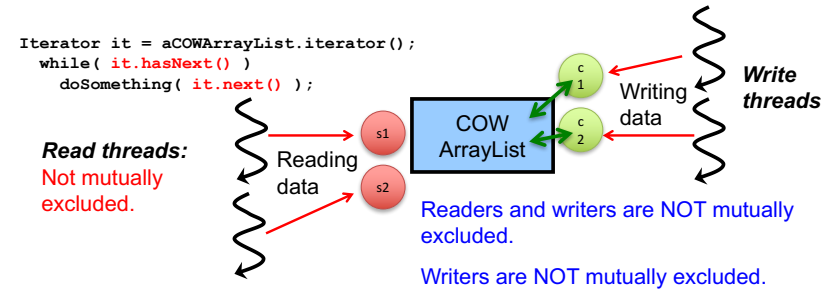
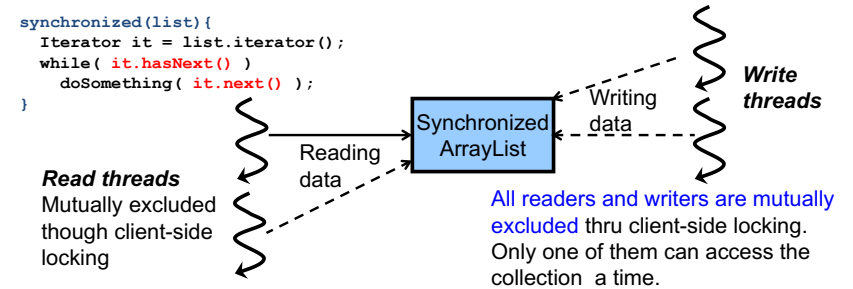


# Copy-On-Write (COW)

- Making **a copy of a collection** when a write thread updates the collection's elements
- A write thread
  - Performs add(), remove(), set() and other state-changing (or mutative) methods on **a duplicate copy** of collection elements.
  - Synchronizes the updated/modified copy with the original element set.
- Write threads ARE NOT mutually excluded **with each other** and **with read threads**.



- Guaranteed that collection elements are never corrupted.
  - Read thread obtains:
    - (10, 0, 1),
    - (10, 0, 1, 300)
    - (10, 0), or
    - (10, 0, 300)
  - It never obtains corrupted list such as (10, 300).



## Pros and Cons in COW Performance

- Pros
  - No concerns about data corruption.
  - Improved performance for iteration
- Cons
  - State-changing methods (e.g., add(), remove()) are very slow.
    - Never use COW collections in single-threaded programs.
    - Their overhead grows exponentially as the number of elements increases.

|   | The overhead of add() [msec] |           |               |              |
|---|------------------------------|-----------|---------------|--------------|
|   | # of elems                   | ArrayList | SyncArrayList | COWArrayList |
| » | 1,000                        | 0         | 0             | 14           |
| » | 5,000                        | 0         | 0             | 102          |
| » | 10,000                       | 0         | 0             | 409          |
| » | 20,000                       | 0         | 0             | 1,712        |
| » | 30,000                       | 15        | 16            | 4,566        |



## When to Use COW Collections?

- Read operations are executed a lot more often than write operations.
- When the # of read threads is a lot greater than the # of write threads.
- When state-changing methods are rarely called.
- When the # of elements is relatively small.

## Lists and Concurrency

- Performance-wise, lists are not that great collections for multi-threaded programs
  - Particularly when a list has many elements.
  - Because state updates are often expensive due to index-based element ordering
    - This makes thread synchronization (locking) more expensive.
- Consider to use other collections
  - Such as queues and sets.
    - Queue: Ordered. Duplicate elements allowed.
    - Set: Not ordered. No duplicated elements allowed.

## In Summary...

- Be aware of the characteristics of COW collections.
- Be conservative to use them.

## Recap: Observable.notifyObservers()

```
• class Observable {  
    private Vector obs;           //observers  
    private boolean changed = false;  
  
    public void notifyObservers(Object arg){  
        Object[] arrLocal;  
        synchronized (this){      // lock.lock();  
            if (!changed) return;  // balking  
            arrLocal = obs.toArray(); // observers copied to arrLocal  
            changed = false;  
        }                          // lock.unlock();  
        for (int i = arrLocal.length-1; i >= 0; i--)  
            ((Observer)arrLocal[i]).update(this, arg); // OPEN CALL  
    }  
    .....  
}
```

## HW 20

- Revise your HW 18 code to replace `ArrayList` with
  - `ConcurrentLinkedQueue`,
  - concurrent hash set (`ConcurrentHashMap.KeySetView<K, Boolean>`),  
or
  - `CopyOnWriteArrayList`

## Observable.notifyObservers()

```
class Observable {
    private ArrayList<Observer> obs;
    private boolean changed = false;
    private ReentrantLock lock;

    void addObserver(Observer o) {
        lock.lock();
        obs.add(o);
        lock.lock();
    }

    void notifyObservers(Object arg) {
        ArrayList<Observer> obsLocal;
        lock.lock();
        if(!changed) return;
        obsLocal =
            new ArrayList<Observer>(obs);
        changed = false;
        lock.unlock();
        Iterator it = obsLocal.iterator();
        while( it.hasNext() ){
            it.next().update(this, arg);
        }
    }
}
```

```
class Observable {
    private CopyOnWriteArrayList<Observer> obs;
    private boolean changed = false;
    private ReentrantLock lock;

    void addObserver(Observer o) {
        obs.add(o);
    }

    void notifyObservers(Object arg) {
        lock.lock();
        if(!changed) return;
        changed = false;
        lock.unlock();
        Iterator it = obs.iterator();
        while( it.hasNext() ){
            it.next().update(this, arg);
        }
    }
}
```