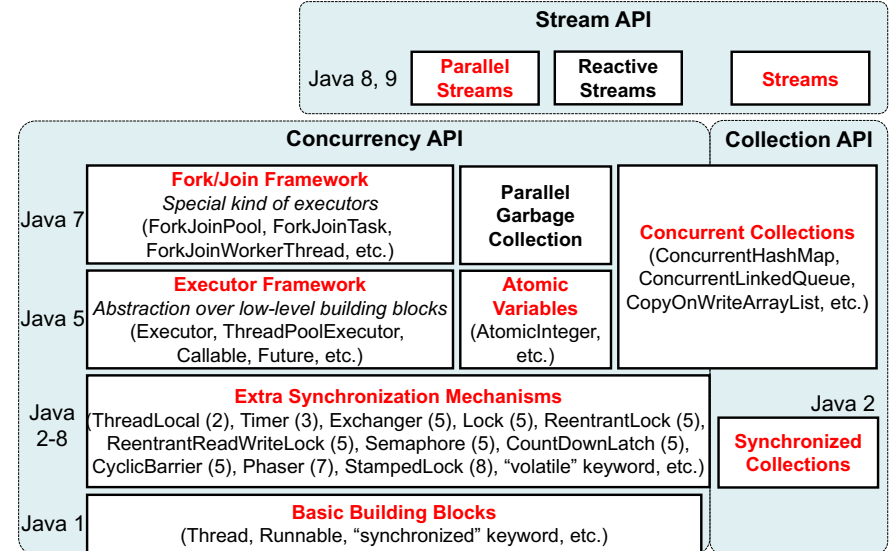


Concurrency API in Java



Executor Framework

Executor Framework


- An abstraction layer atop low-level concurrency primitives
 - Focuses on task execution on threads
 - Decouples task execution (on threads) from task submission (to threads) to make task execution configurable.
 - Introduced in Java 5 (2004)
 - Further enhanced in subsequent versions
 - Implemented in `java.util.concurrent`.

Tasks, Threads and Executor

- Tasks
 - Logical units of work
 - e.g., prime number generation, file caching, file crawling, file indexing, access counting
- Threads
 - Mechanism to run tasks *concurrently*.
- Executor
 - Is the primary abstraction for task execution
 - `Thread` is not anymore.

Executor

- ```
public interface Executor{
 void execute(Runnable command);
}
```
- `Runnable`'s `run()` implements a task.
- Follows the Producer-Consumer design pattern.
  - Producers: submit tasks
  - Consumers: execute tasks
- Provides the easiest way to implement producers and consumers
- Makes task execution configurable.



5


## Task Execution Policies

- The Executor framework allows you to specify and customize the *execution policy* for tasks.
  - “What, where, when and how” of task execution.
    - In what thread will tasks be executed?
    - In what order should tasks be executed (FIFO, LIFO, priority-based ordering)?
    - How many tasks may execute concurrently.
    - How many tasks may be queued pending execution?
    - If a task has to be rejected because an application is overloaded, which task should be selected as the victim? How should the application be notified?
    - What actions should be taken before or after executing a task?

6


## Thread Pool

- A key component for task execution.
- A set of *pre-created “worker” threads* that will be used for future task execution
- Each worker thread
  - Gets and executes a task if it is available in the queue.
  - Goes to the Waiting state until a producer submits the next task.



7

- Benefits of using a thread pool
  - Can eliminate runtime overhead to create threads
  - Can bound the maximum number of threads (i.e., the max amount of resource utilization)



8

# Executors

- A utility class for `Executor` objects
- Defines static factory methods to create thread pools.

```
- static ExecutorService newFixedThreadPool(int n)
 • Fixed-size thread pool.

- ExecutorService executor = Executors.newFixedThreadPool(2);
 executor.execute(new PrimeNumberGenerator(1L, 500000L));
 executor.execute(new PrimeNumberGenerator(500001L, 1000000L));

- Thread t1, t2;
 t1 = new Thread(new PrimeNumberGenerator(1L, 500000L));
 t2 = new Thread(new PrimeNumberGenerator(500001L, 1000000L));
 t1.start();
 t2.start();
```

9

- Defines static factory methods to create thread pools.

```
- static ExecutorService newFixedThreadPool(int n)
 • Fixed-size thread pool.

- static ScheduledExecutorService newScheduledThreadPool(int n)
 • Fixed-size thread pool that supports delayed and periodic task execution.


- static ExecutorService newSingleThreadExecutor()
 • A pool that operates only one thread.

- static ScheduledExecutorService newSingleThreadScheduledExecutor()
 • A single-threaded pool that supports delayed and periodic task execution.
```

10

- ```
- static ExecutorService newCachedThreadPool()
    • Variable-size (not fixed-size) thread pool.
      - Uses previously created "idle" threads if they are available.
      - Creates a new thread if no idle threads are available.
      - Idle threads are terminated and removed from the pool after they are not used for 60 seconds.

    • Pros: Intends to minimize the number of tasks in the queue.
    • Cons: No cap for the number of threads in the pool.
    • Useful to handle a number of short-lived (lightweight) tasks
```




11

- ```
- static ExecutorService newWorkStealingPool(int parallelism)
 • Variable-size (not fixed-size) thread pool with a cap for the number of threads.
 - parallelism specifies the cap for the number of threads.

 • Each worker thread
 - Has its own "primary" queue and gets the next task from the queue.
 - "Steals" a task from another queue if no tasks are available in its primary queue.
 - Dies after being idle for some time.


 • Pros:
 - Each queue requires less thread synchronization.
 - Each queue operates to minimize the number of tasks.
 - The number of worker threads can be bounded.

 • Cons: N/A
```



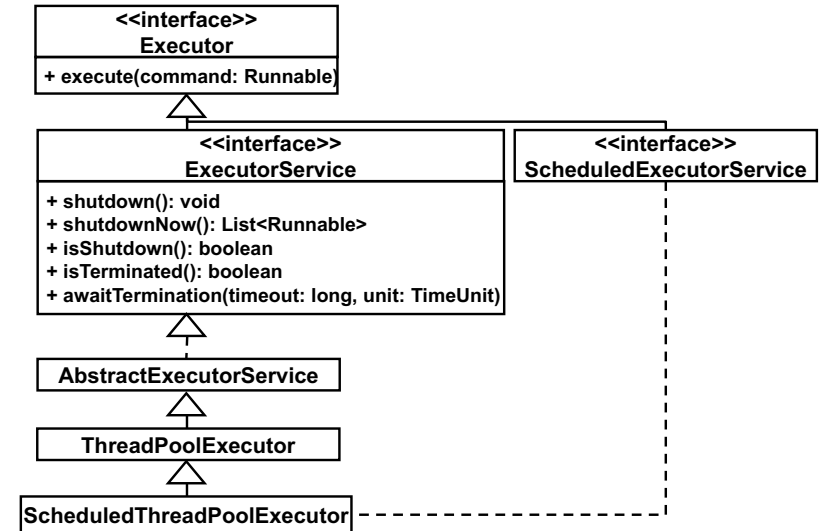
12

- `static ExecutorService newWorkStealingPool()`
  - Obtains the number of available CPU cores by calling `availableProcessors()` and invokes the previous version of `newWorkStealingPool()`



13

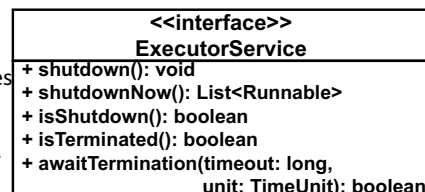
## ExecutorService




14

## Termination of Executor

- Methods to terminate an executor
  - `shutdown()`
    - Rejects new tasks to get in
      - Throws a `RejectedExecutionException`
    - Allows previously submitted tasks to complete
      - Tasks being executed
      - Tasks in the queue
  - `shutdownNow()`
    - Rejects new tasks to get in
    - Removes all tasks from the queue and returns them
    - Tries to stop the tasks that are being executed.
      - Call `interrupt()` on each worker thread
    - A task can be stopped if it checks `Thread.interrupted()` or catches `InterruptedException` to exit `run()`.
      - Otherwise, it may not be stopped.



- 3 states of an executor
  - **Running**
  - **Shutting down**
    - Once `shutdown()` or `shutdownNow()` is called.
    - `isShutdown()` returns true.
  - **Terminated**
    - Once all tasks have been completed or stopped.
    - `isTerminated()` returns true.



- Use `awaitTermination()` if you want to wait for an executor to be terminated.

- It blocks until the pool is terminated or the timeout occurs.
- It returns true if the pool is terminated or false otherwise.

```
- executor.shutdown();
 executor.awaitTermination(Long.MAX_VALUE, TimeUnit.SECONDS);
 doSomething();

- executor.shutdown();
 if(!executor.awaitTermination(60, TimeUnit.SECONDS)){
 shutdownNow();
 }
 if(!executor.awaitTermination(60, TimeUnit.SECONDS)){
 doErrorHandling();
 }
```

| <<interface>><br>ExecutorService                              |  |
|---------------------------------------------------------------|--|
| + shutdown(): void                                            |  |
| + shutdownNow(): List<Runnable>                               |  |
| + isShutdown(): boolean                                       |  |
| + isTerminated(): boolean                                     |  |
| + awaitTermination(timeout: long,<br>unit: TimeUnit): boolean |  |

```
• Runnable r1, r2;
 r1 = new PrimeNumberGenerator(1L, 500000L);
 r2 = new PrimeNumberGenerator(500001L, 1000000L);
 ExecutorService executor = Executors.newFixedThreadPool(2);
 executor.execute(r1);
 executor.execute(r2);
 executor.shutdown();
 executor.awaitTermination(...);
 r1.getPrimes().forEach(...);
 r2.getPrimes().forEach(...);
```